

UNM SCHOOL of ENGINEERING
Department of Computer Science

Information Storage, Integer Representation and Arithmetic

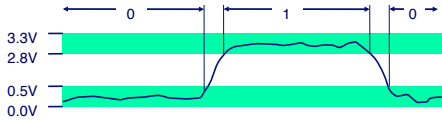
CS 341: Intro. to Computer Architecture & Organization

Andree Jacobson
(Slides by Prof. Dorian Arnold)

Today: Bits, Bytes, and Integers

- ▶ Representing information as bits
- ▶ Bit-level manipulations
- ▶ Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- ▶ Summary

Binary Representations



- ▶ Three most important binary number representations
 - Unsigned integers
 - Two's complement integers (for signed integers)
 - Floating point numbers

Encoding Byte Values

► Byte = 8 bits

- Binary 00000000₂ to 11111111₂
- Decimal: 0₁₀ to 255₁₀
- Hexadecimal 00₁₆ to FF₁₆
 - Base 16 number representation
 - Use characters '0' to '9' and 'A' to 'F'
 - Write FA1D37B₁₆ in C as
 - 0xFA1D37B

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Byte-Oriented Memory Organization



- Programs reference (virtual) addresses for data/instr.
 - Simple concept: very large byte array
 - Implemented with hierarchy of different memory types
 - System provides address space private to particular "process"
 - Program being executed
 - Program can clobber its own data, but not that of others
- Compiler + Run-Time System Control Allocation
 - Where different program objects should be stored
 - All allocation within single virtual address space

Machine Words


- Smallest unit addressable by machine
 - Nominal size of integer-valued data
 - Most current machines (still) use 32 bits
 - 4-byte word/pointer limits address spaces to 4GB
 - High-end systems use 64 bits (8 bytes) words
 - Potential address space $\approx 1.8 \times 10^{19}$ bytes
 - x86-64 machines support 48-bit addresses: 256 Terabytes
 - Machines can support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

Data Representations

C Data Type	Typical 32-bit	Intel IA32	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	4	8
long long	8	8	8
float	4	4	4
double	8	8	8
long double	8	10/12	10/16
pointer	4	4	8

Referencing (multi-byte) Data

- Address specifies location of first byte
- Treatment of addresssed byte(s) depends on reference type




Byte Ordering

- How should bytes within a multi-byte word be ordered in memory?
- Conventions
 - Big Endian: Sun, PPC Mac, Internet
 - Least significant byte has highest address
 - Little Endian: x86
 - Least significant byte has lowest address

Byte Ordering Example

- ▶ Big Endian
 - Least significant byte has highest address
- ▶ Little Endian
 - Least significant byte has lowest address
- ▶ Example
 - Variable x has 4-byte representation 0x01234567
 - Address given by &x is 0x100



Examining Data Representations

- ▶ Code to Print Byte Representation of Data
 - Casting pointer to unsigned char * creates byte array

```
typedef unsigned char *pointer;
void show_bytes(pointer start, int len) {
    int i;
    for (i = 0; i < len; i++)
        printf("%p\t0x%.2x\n", start+i, start[i]);
}
```

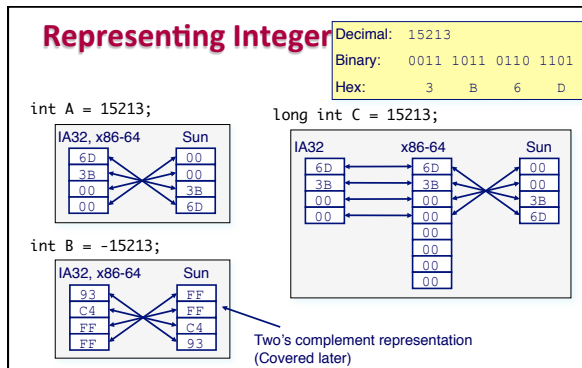
Printf directives:
 %p: Print pointer
 %x: Print Hexadecimal

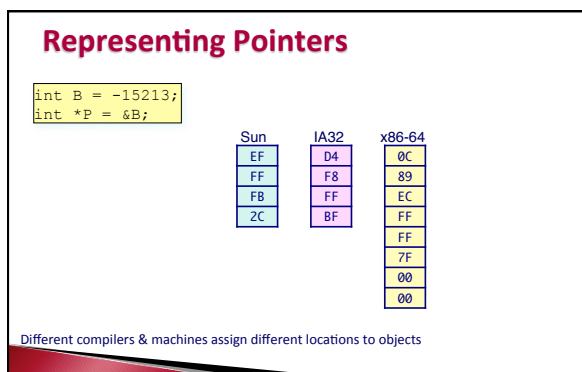
show_bytes Execution Example

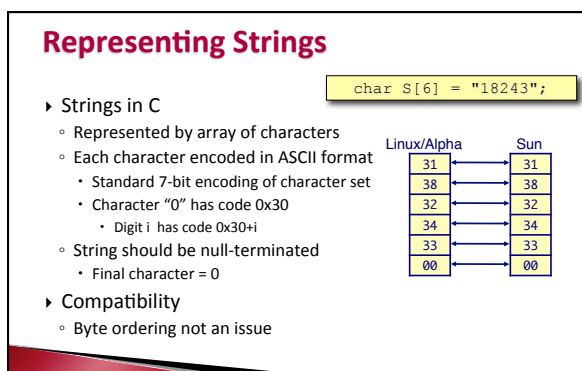
```
int a = 15213;
printf("int a = 15213;\n");
show_bytes((pointer) &a, sizeof(int));
```

Result (Linux):

```
int a = 15213;
0x11ffffcb8 0x6d
0x11ffffcb9 0x3b
0x11ffffcba 0x00
0x11ffffcbb 0x00
```







Data Representation Example

- ▶ Is it more compact to represent positive numerical data as an integer or a string of characters?
 - Ignore arithmetic operations, just consider data size!
 - It depends!
 - On the range of numbers to be represented.
 - Integers will always need 4 (or 8) bytes per number.
 - String of chars need 1 byte per digit (plus null-terminator)
 - More compact for numbers from 0-999

Today: Bits, Bytes, and Integers

- ▶ Representing information as bits
- ▶ Bit-level manipulations
- ▶ Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- ▶ Summary

Boolean Algebra

- ▶ Developed by George Boole in 19th Century
 - Algebraic representation of logic
 - Encode "True" as 1 and "False" as 0
- And

■ $A \& B = 1$ when both $A=1$ and $B=1$

	A	B
0	0	1
0	0	0
1	0	1

Not

■ $\sim A = 1$ when $A=0$

	A
0	1
1	0

Or

■ $A | B = 1$ when either $A=1$ or $B=1$

	A	B
0	0	1
0	0	1
1	1	1

Exclusive-Or (Xor)

■ $A \wedge B = 1$ when either $A=1$ or $B=1$, but not both

	A	B
0	0	1
0	0	1
1	1	0

General Boolean Algebras

► Operate on Bit Vectors

- Operations applied bitwise

01101001	01101001	01101001
<u>& 01010101</u>	<u> 01010101</u>	<u>^ 01010101</u>
01000001	01111101	00111100
		<u>~ 01010101</u>
		10101010

► All of the Properties of Boolean Algebra Apply

Representing & Manipulating Sets

► Representation

- Width w bit vector represents subsets of {0, ..., w-1}
- $a_j = 1$ if $j \in A$

- 01101001 { 0, 3, 5, 6 }
- 76543210

- 01010101 { 0, 2, 4, 6 }
- 76543210

► Operations

- | | | | |
|-----|----------------------|----------|----------------------|
| ◦ & | Intersection | 01000001 | { 0, 6 } |
| ◦ | Union | 01111101 | { 0, 2, 3, 4, 5, 6 } |
| ◦ ^ | Symmetric difference | 00111100 | { 2, 3, 4, 5 } |
| ◦ ~ | Complement | 10101010 | { 1, 3, 5, 7 } |

Bit-Level Operations in C

► Operations &, |, ~, ^ available in C

- Apply to any "integral" data type
 - long, int, short, char, unsigned
- View arguments as bit vectors
- Operations applied bit-wise

► Examples (Char data type)

- ~0x41 → 0xBE
- ~01000001₂ → 10111110₂
- ~0x00 → 0xFF
- ~00000000₂ → 11111111₂
- 0x69 & 0x55 → 0x41
- 01101001₂ & 01010101₂ → 01000001₂
- 0x69 | 0x55 → 0x7D
- 01101001₂ | 01010101₂ → 01111101₂

Contrast: Logic Operations in C

- ▶ Contrast to Logical Operators
 - `&&`, `||`, `!`
 - `0` → "False"; nonzeros → "True"
 - Always return 0 or 1 (often interpreted as true or false)
 - **Early termination**
- ▶ Examples (char data type)
 - `!0x41` → `0x00`
 - `!0x00` → `0x01`
 - `!!0x41` → `0x01`
 - `0x69 && 0x55` → `0x01`
 - `0x69 || 0x55` → `0x01`

Shift Operations

- ▶ Left Shift: `x << y`
 - Shift bit-vector `x` left `y` positions
 - Throw away extra bits on left
 - Fill with 0's on right
- ▶ Right Shift: `x >> y`
 - Shift bit-vector `x` right `y` positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on right
- ▶ Undefined Behavior
 - Shift amount `< 0` or `≥` word size

Argument <code>x</code>	01100010
<code><< 3</code>	00010000
Log. <code>>> 2</code>	00011000
Arith. <code>>> 2</code>	00011000

Argument <code>x</code>	10100010
<code><< 3</code>	00010000
Log. <code>>> 2</code>	00101000
Arith. <code>>> 2</code>	11101000

Today: Bits, Bytes, and Integers

- ▶ Representing information as bits
- ▶ Bit-level manipulations
- ▶ Integers
 - **Representation: unsigned and signed**
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- ▶ Summary

Encoding Integers

Unsigned

$$B2U_w(X) = \sum_{i=0}^{w-1} x_i 2^i$$

Two's Complement

$$B2T_w(X) = -x_{w-1} 2^{w-1} + \sum_{i=0}^{w-2} x_i 2^i$$

```
short int x = 15213;
short int y = -15213;
```

Sign Bit

Sign Bit

- For 2's complement, most significant bit indicates sign
 - 0 for nonnegative
 - 1 for negative

Encoding Example (Cont.)

```
x = 15213: 00111011 01101101
y = -15213: 11000100 10010011
```

Weight	15213	-15213
1	1	1
2	0	1
4	1	0
8	1	0
16	0	1
32	1	0
64	1	0
128	0	1
256	1	0
512	1	0
1024	0	1
2048	1	0
4096	1	0
8192	1	0
16384	0	1
-32768	0	1
Sum	15213	-15213

Numeric Ranges

Unsigned Values

- $UMin = 0$
000...0
- $UMax = 2^w - 1$
111...1

Two's Complement Values

- $TMin = -2^{w-1}$
100...0
- $TMax = 2^{w-1} - 1$
011...1

Other Values

- Minus 1
111...1

Values for $W = 16$

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Values for Different Word Sizes

	W			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

Observations

- $|TMin| = TMax + 1$
 - Asymmetric range
- $UMax = 2 * TMax + 1$

C Programming

- `#include <limits.h>`
- Declares constants, e.g.,
 - `ULONG_MAX`
 - `LONG_MAX`
 - `LONG_MIN`
- Values platform specific

Unsigned & Signed Numeric Values

X	B2U(X)	B2T(X)
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

Equivalence

- Same encodings for nonnegative values

Uniqueness

- Every bit pattern represents unique integer value
- Each representable integer has unique bit encoding


⇒ Can Invert Mappings

- $U2B(x) = B2U^{-1}(x)$
 - Bit pattern for unsigned integer
- $T2B(x) = B2T^{-1}(x)$
 - Bit pattern for two's comp integer

Today: Bits, Bytes, and Integers


- ▶ Representing information as bits
- ▶ Bit-level manipulations
- ▶ Integers
 - Representation: unsigned and signed
 - **Conversion, casting**
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- ▶ Summary

Mapping Between Signed & Unsigned



- Mappings between unsigned and two's complement:
keep bit representations and reinterpret


Relation between Signed & Unsigned

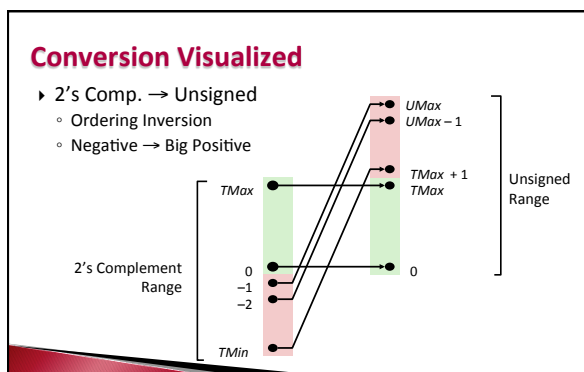


Mapping Signed ↔ Unsigned

Bits	Signed	Unsigned
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	-8	8
1001	-7	9
1010	-6	10
1011	-5	11
1100	-4	12
1101	-3	13
1110	-2	14
1111	-1	15

Mapping: Signed \leftrightarrow Unsigned (offset of ± 8 for signed, ± 16 for unsigned)





Signed vs. Unsigned in C

- Constants
 - By default are considered to be signed integers
 - Unsigned if have "U" as suffix
0U, 4294967259U
- Casting
 - Explicit casting between signed & unsigned same as U2T and T2U


```
int tx, ty;
unsigned ux, uy;
tx = (int) ux;
uy = (unsigned) ty;
```
 - Implicit casting also occurs via assignments and procedure calls


```
tx = ux;
uy = ty;
```

Casting Surprises

- ▶ Expression Evaluation
 - If there is a mix of unsigned and signed in single expression, **signed values implicitly cast to unsigned**
 - Including comparison operations <, >, ==, <=, >=
 - Examples for W = 32:
TMIN = -2,147,483,648, TMAX = 2,147,483,647

Constant ₁	Constant ₂	Relation	Eval.
0	0U	==	unsigned
-1	0	<	signed
-1	0U	>	unsigned
2147483647	-2147483647-1	<	signed
2147483647U	-2147483647-1	<	unsigned
-1	-2	>	signed
(unsigned)-1	-2	>	unsigned
2147483647	2147483648U	<	unsigned
2147483647	(int) 2147483648U	>	signed

Code Security Example

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

- ▶ Similar code found in FreeBSD's getpeername()
- ▶ There are legions of smart people trying to find vulnerabilities in programs

Typical Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```

Malicious Usage

```

/* Declaration of library function memcpy */
#define size_t unsigned
void *memcpy(void *dest, void *src, size_t n);

/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}

#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    . . .
}

```

Summary

Casting Signed ↔ Unsigned: Basic Rules


- ▶ Bit pattern is maintained
- ▶ But reinterpreted
- ▶ Can have unexpected effects: adding or subtracting 2^w
- ▶ Expression containing signed and unsigned int
 - int is cast to unsigned!!

Today: Bits, Bytes, and Integers

- ▶ Representing information as bits
- ▶ Bit-level manipulations
- ▶ Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - **Expanding, truncating**
 - Addition, negation, multiplication, shifting
- ▶ Summary

Sign Extension

- ▶ Task:
 - Given w -bit signed integer x
 - Convert it to $w+k$ -bit integer with same value
- ▶ Rule:
 - Make k copies of sign bit:
 - $X' = X_{w-1}, \dots, X_{w-1}, X_{w-1}, X_{w-2}, \dots, X_0$



Sign Extension Example

```
short int x = 15213;
int ix = (int) x;
short int y = -15213;
int iy = (int) y;
```

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
ix	15213	00 00 3B 6D	00000000 00000000 00111011 01101101
y	-15213	C4 93	11000100 10010011
iy	-15213	FF FF C4 93	11111111 11111111 11000100 10010011

- ▶ Converting from smaller to larger integer data type
- ▶ C automatically performs sign extension

Summary:

Expanding, Truncating: Basic Rules

- ▶ Expanding (e.g., short int to int)
 - Unsigned: zeros added
 - Signed: sign extension
 - Both yield expected result
- ▶ Truncating (e.g., unsigned to unsigned short)
 - Unsigned/signed: bits are truncated
 - Result reinterpreted
 - Unsigned: mod operation
 - Signed: similar to mod
 - For small numbers yields expected behaviour

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - **Addition, negation, multiplication, shifting**
- Summary

Negation: Complement & Increment

- Claim: Following Holds for 2's Complement

$$\sim x + 1 == -x$$

- Complement

- Observation: $\sim x + x == 1111.111 == -1$

$$\begin{array}{r} x \quad 10011101 \\ + \sim x \quad 01100010 \\ \hline -1 \quad 11111111 \end{array}$$

- Complete Proof?

Complement & Increment Examples

x = 15213

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
~x	-15214	C4 92	11000100 10010010
~x+1	-15213	C4 93	11000100 10010011
y	-15213	C4 93	11000100 10010011

x = 0


	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
~0	-1	FF FF	11111111 11111111
~0+1	0	00 00	00000000 00000000

Unsigned Addition

Operands: w bits

True Sum: $w+1$ bits

Discard Carry: w bits




- ▶ Standard Addition Function
 - Ignores carry output
- ▶ Implements Modular Arithmetic

$$s = UAdd_w(u, v) = (u + v) \% 2^w$$

$$UAdd_w(u, v) = \begin{cases} u + v & u + v < 2^w \\ u + v - 2^w & u + v \geq 2^w \end{cases}$$



Visualizing (Mathematical) Integer Addition

- ▶ Integer Addition
 - 4-bit integers u, v
 - Compute true sum $Add_8(u, v)$
 - Values increase linearly with u and v
 - Forms planar surface



Visualizing Unsigned Addition

- ▶ Wraps Around
 - If true sum $\geq 2^w$
 - At most once



Mathematical Properties

► Modular Addition Forms an *Abelian Group*

- **Closed** under addition
 $0 \leq \text{UAdd}_w(u, v) \leq 2^w - 1$
- **Commutative**
 $\text{UAdd}_w(u, v) = \text{UAdd}_w(v, u)$
- **Associative**
 $\text{UAdd}_w(t, \text{UAdd}_w(u, v)) = \text{UAdd}_w(\text{UAdd}_w(t, u), v)$
- **0** is additive identity
 $\text{UAdd}_w(u, 0) = u$
- Every element has additive **inverse**
 - Let $\text{UComp}_w(u) = 2^w - u$
 $\text{UAdd}_w(u, \text{UComp}_w(u)) = 0$


Two's Complement Addition

Operands: w bits



True Sum: $w+1$ bits

Discard Carry: w bits



► TAdd and UAdd have Identical Bit-Level Behavior


- Signed vs. unsigned addition in C:

```
int s, t, u, v;
s = (int) ((unsigned) u + (unsigned) v);
t = u + v;
// Will give s == t
```

TAdd Overflow


► Functionality

- True sum requires $w+1$ bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer



Visualizing 2's Complement Addition

- ▶ Values
 - 4-bit two's comp.
 - Range from -8 to +7
- ▶ Wraps Around
 - If $\text{sum} \geq 2^{w-1}$
 - Becomes negative
 - At most once
 - If $\text{sum} < -2^{w-1}$
 - Becomes positive
 - At most once



Mathematical Properties of TAdd

- ▶ Isomorphic Group to unsigneds with UAdd
 - $TAdd_w(u, v) = U2T(UAdd_w(T2U(u), T2U(v)))$
 - Since both have identical bit patterns
- ▶ Two's Complement Under TAdd Forms a Group
 - Closed, Commutative, Associative, 0 is additive identity
 - Every element has additive inverse

$$TComp_w(u) = \begin{cases} -u & u \neq TMin_w \\ TMin_w & u = TMin_w \end{cases}$$

Multiplication


- ▶ Computing Exact Product of w -bit numbers x, y
 - Either signed or unsigned
- ▶ Ranges
 - Unsigned: $0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$
 - Up to $2w$ bits
 - Two's complement min: $x * y \geq (-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$
 - Up to $2w-1$ bits
 - Two's complement max: $x * y \leq (2^{w-1} - 1)^2 = 2^{2w-2} - 2^{w-1} + 1$
 - Up to $2w$ bits, but only for $(TMin_w)^2$
- ▶ Maintaining Exact Results
 - Would need to keep expanding word size with each product computed
 - Done in software by "arbitrary precision" arithmetic packages

Unsigned Multiplication in C

Operands: w bits

True Product: $2w$ bits

Discard w bits: w bits



Standard Multiplication Function

- Ignores high order w bits

Implements Modular Arithmetic

$$UMult_w(u, v) = u \cdot v \bmod 2^w$$

Code Security Example #2: XDR Code

```
void* copy_elements(void *ele_src[], int ele_cnt, size_t ele_size) {
    /*
     * Allocate buffer for ele_cnt objects, each of ele_size bytes
     * and copy from locations designated by ele_src
     */
    void *result = malloc(ele_cnt * ele_size);
    if (result == NULL)
        /* malloc failed */
        return NULL;
    void *next = result;
    int i;
    for (i = 0; i < ele_cnt; i++) {
        /* Copy object i to destination */
        memcpy(next, ele_src[i], ele_size);
        /* Move pointer to next memory region */
        next += ele_size;
    }
    return result;
}
```

XDR Vulnerability

`malloc(ele_cnt * ele_size)`

What if:

- **ele_cnt** = $2^{20} + 1$
- **ele_size** = 4096 = 2^{12}
- Allocation = ??


How can I make this function secure?

Signed Multiplication in C

Operands: w bits

True Product: $2w$ bits

Discard w bits: w bits



Standard Multiplication Function

- Ignores high order w bits
- Some of which are different for signed vs. unsigned multiplication
- Lower bits are the same

Power-of-2 Multiply with Shift


Operation

- $u \ll k$ gives $u \cdot 2^k$
- Both signed and unsigned

Operands: w bits

True Product: $w+k$ bits

Discard k bits: w bits



Examples

- $u \ll 3 \quad == \quad u \cdot 8$
- $u \ll 5 - u \ll 3 \quad == \quad u \cdot 24$
- Most machines shift and add faster than multiply
 - Compiler generates this code automatically

Compiled Multiplication Code

C Function

```
int mul12(int x)
{
    return x*12;
}
```

Compiled Arithmetic Operations

```
leal (%eax,%eax,2), %eax
sall $2, %eax
```

Explanation

```
t <- x+x*2
return t << 2;
```


- C compiler automatically generates shift/add code when multiplying by constant

Unsigned Power-of-2 Divide with Shift

Quotient of Unsigned by Power of 2

- $u \gg k$ gives $\lfloor u / 2^k \rfloor$

- Uses logical shift



	Division	Computed	Hex	Binary
x	15213	15213	3B 6D	00111011 01101101
$x \gg 1$	7606.5	7606	1D B6	00011101 10110110
$x \gg 4$	950.8125	950	03 B6	00000011 10110110
$x \gg 8$	59.4257813	59	00 3B	00000000 00111011

Compiled Unsigned Division Code

C Function

```
unsigned udiv8(unsigned x)
{
    return x/8;
}
```

Compiled Arithmetic Operations

```
shrl $3, %eax
```

Explanation

```
# Logical shift
return x >> 3;
```

- Uses logical shift for unsigned
- For Java Users
 - Logical shift written as \gg


Signed Power-of-2 Divide with Shift

Quotient of Signed by Power of 2

- $x \gg k$ gives $\lfloor x / 2^k \rfloor$

- Uses arithmetic shift

- Rounds wrong direction when $u < 0$




	Division	Computed	Hex	Binary
y	-15213	-15213	C4 93	11000100 10010011
$y \gg 1$	-7606.5	-7607	E2 49	11100010 01001001
$y \gg 4$	-950.8125	-951	FC 49	11111100 01001001
$y \gg 8$	-59.4257813	-60	FF C4	11111111 11000100

Correct Power-of-2 Divide


- Quotient of Negative Number by Power of 2
 - Want $\lfloor x / 2^k \rfloor$ (Round Toward 0)
 - Bias dividend toward 0 (Additive Bias: $2^k - 1$)
 - Compute as $\lfloor (x + 2^k - 1) / 2^k \rfloor$
 - In C: $(x + (1 < k) - 1) >> k$

Case 1: No rounding



Correct Power-of-2 Divide (Cont.)

Case 2: Rounding



Compiled Signed Division Code

C Function

```
int idiv8(int x)
{
    return x/8;
}
```

Compiled Arithmetic Operations

```
testl %eax, %eax
js L4
L3:
    sarl $3, %eax
    ret
L4:
    addl $7, %eax
    jmp L3
```

Explanation

```
if x < 0
    x += 7;
# Arithmetic shift
return x >> 3;
```

- Uses arithmetic shift for int
- For Java Users
 - Arith. shift written as \gg

Arithmetic: Basic Rules

- ▶ Addition:
 - Unsigned/signed: Normal addition followed by truncate, same operation on bit level
 - Unsigned: addition mod 2^w
 - Mathematical addition + possible subtraction of 2^w
 - Signed: modified addition mod 2^w (result in proper range)
 - Mathematical addition + possible addition or subtraction of 2^w
- ▶ Multiplication:
 - Unsigned/signed: Normal multiplication followed by truncate, same operation on bit level
 - Unsigned: multiplication mod 2^w
 - Signed: modified multiplication mod 2^w (result in proper range)

Arithmetic: Basic Rules

- ▶ Unsigned ints, 2's complement ints are isomorphic rings: isomorphism = casting
- ▶ Left shift
 - Unsigned/signed: multiplication by 2^k
 - Always logical shift
- ▶ Right shift
 - Unsigned: logical shift, div (division + round to zero) by 2^k
 - Signed: arithmetic shift
 - Positive numbers: div (division + round to zero) by 2^k
 - Negative numbers: div (division + round away from zero) by 2^k

Today: Integers

- ▶ Representation: unsigned and signed
- ▶ Conversion, casting
- ▶ Expanding, truncating
- ▶ Addition, negation, multiplication, shifting
- ▶ Summary

Properties of Unsigned Arithmetic

- ▶ Unsigned Multiplication with Addition Forms Commutative Ring
 - Addition is commutative group
 - Closed under multiplication
 - $0 \leq \text{UMult}_w(u, v) \leq 2^w - 1$
 - Multiplication Commutative
 - $\text{UMult}_w(u, v) = \text{UMult}_w(v, u)$
 - Multiplication is Associative
 - $\text{UMult}_w(t, \text{UMult}_w(u, v)) = \text{UMult}_w(\text{UMult}_w(t, u), v)$
 - 1 is multiplicative identity
 - $\text{UMult}_w(u, 1) = u$
 - Multiplication distributes over addition
 - $\text{UMult}_w(t, \text{UAdd}_w(u, v)) = \text{UAdd}_w(\text{UMult}_w(t, u), \text{UMult}_w(t, v))$

Properties of Two's Comp. Arithmetic

- ▶ Isomorphic Algebras
 - Unsigned multiplication and addition
 - Truncating to w bits
 - Two's complement multiplication and addition
 - Truncating to w bits
- ▶ Both Form Rings
 - Isomorphic to ring of integers mod 2^w
- ▶ Comparison to (Mathematical) Integer Arithmetic
 - Both are rings
 - Integers obey ordering properties, e.g.,
 - $u > 0 \implies u + v > v$
 - $u > 0, v > 0 \implies u \cdot v > 0$
 - These properties are not obeyed by two's comp. arithmetic
 - $TMax + 1 == TMin$
 - $15213 * 30426 == -10030$ (16-bit words)

Why Should I Use Unsigned?

- ▶ *Don't* Use Just Because Number Nonnegative
 - Easy to make mistakes


```
unsigned i;
for (i = cnt-2; i >= 0; i--)
  a[i] += a[i+1];
```
 - Can be very subtle


```
#define DELTA sizeof(int)
int i;
for (i = CNT; i-DELTA >= 0; i-= DELTA)
  . . .
```
- ▶ *Do* Use When Performing Modular Arithmetic
 - Multiprecision arithmetic
- ▶ *Do* Use When Using Bits to Represent Sets
 - Logical right shift, no sign extension

Integer C Puzzles

Initialization

```
int x = foo();
int y = bar();
unsigned ux = x;
unsigned uy = y;
```

- $x < 0 \Rightarrow ((x*2) < 0)$
- $ux \geq 0$
- $x \& 7 == 7 \Rightarrow (x < 30) < 0$
- $ux > -1$
- $x > y \Rightarrow -x < -y$
- $x * x \geq 0$
- $x > 0 \&\& y > 0 \Rightarrow x + y > 0$
- $x \geq 0 \Rightarrow -x \leq 0$
- $x \leq 0 \Rightarrow -x \geq 0$
- $(x|-x) >> 31 == -1$
- $ux >> 3 == ux/8$
- $x >> 3 == x/8$
