

Floating Point

CS 341: Intro. to Computer
Architecture & Organization

Andree Jacobson
(Slides courtesy of
Prof. Dorian Arnold)


Today: Floating Point

- ▶ Background: Fractional binary numbers
- ▶ IEEE floating point standard: Definition
- ▶ Example and properties
- ▶ Rounding, addition, multiplication
- ▶ Floating point in C
- ▶ Summary

Fractional binary numbers

- ▶ What is 1011.101_2 ?

Fractional Binary Numbers



- ▶ Representation
 - Bits to right of “binary point” represent fractional powers of 2
 - Represents rational number: $\sum_{k=-j}^i b_k \times 2^k$

Fractional Binary Numbers: Examples

- | Value | Representation |
|-------|-----------------------|
| 5 3/4 | 101.11 ₂ |
| 2 7/8 | 10.111 ₂ |
| 63/64 | 0.111111 ₂ |
- Observations
 - Divide by 2 by shifting right
 - Multiply by 2 by shifting left
 - Numbers of form 0.111111...₂ are just below 1.0
 - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
 - Use notation $1.0 - \epsilon$

Representable Numbers

- ▶ Limitation
 - Can only exactly represent numbers of the form $x2^k$
 - Other rational numbers have repeating bit representations
- ▶ Value Representation
 - 1/3 0.0101010101[01]...₂
 - 1/5 0.001100110011[0011]...₂
 - 1/10 0.0001100110011[0011]...₂

Today: Floating Point

- ▶ Background: Fractional binary numbers
- ▶ IEEE floating point standard: Definition
- ▶ Example and properties
- ▶ Rounding, addition, multiplication
- ▶ Floating point in C
- ▶ Summary

IEEE Floating Point

- ▶ IEEE Standard 754
 - Established in 1985 as uniform standard for floating point arithmetic
 - Before that, many idiosyncratic formats
 - Supported by all major CPUs
- ▶ Driven by numerical concerns
 - Nice standards for rounding, overflow, underflow
 - Hard to make fast in hardware
 - Numerical analysts predominated over hardware designers in defining standard

Floating Point Representation

► Numerical Form:

$$(-1)^s M 2^E$$

- Sign bit **s** determines whether number is negative or positive
- Significand **M** normally a fractional value in range [1.0,2.0).
- Exponent **E** weights value by power of two


► Encoding

- MSB **S** is sign bit **s**
- **exp** field encodes **E** (but is not equal to E)
- **frac** field encodes **M** (but is not equal to M)




Precisions


► Single precision: 32 bits



► Double precision: 64 bits



► Extended precision: 80 bits (Intel only)



Normalized Values

► Condition: exp ≠ 000...0 and exp ≠ 111...1

► Exponent coded as biased value: E = Exp – Bias

- Exp: unsigned value exp
- Bias = $2^{k-1} - 1$, where k is number of exponent bits
 - Single precision: Bias: 127 (Exp: 1...254, E: -126...127)
 - Double precision: Bias: 1023 (Exp: 1...2046, E: -1022...1023)

► Significand coded with implied leading 1: M = 1.XXX...X₂

- XXX...X: bits of **frac**
- Minimum when 000...0 (M = 1.0)
- Maximum when 111...1 (M = 2.0 – ε)
- Get extra leading bit for “free”

Normalized Encoding Example

► Value: Float F = 15213.0;

$$15213_{10} = 11101101101101_2 = 1.1101101101101_2 \times 2^{13}$$

► Significand

$$M = 1.1101101101101_2$$

$$\text{frac} = 11011011011010000000000_2$$


► Exponent

$$E = 13$$

$$\text{Bias} = 127$$

$$\text{Exp} = 140 = 10001100_2$$

► Result:




Denormalized Values

- Condition: $\text{exp} = 000\dots 0$
- Exponent value: $E = 1 - \text{Bias}$ (instead of $E = 0 - \text{Bias}$)
- Significand coded with implied leading 0: $M = 0.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}_2$: bits of frac
- Cases
 - $\text{exp} = 000\dots 0, \text{frac} = 000\dots 0$
 - Represents zero value
 - Note distinct values: $+0$ and -0 (why?)
 - $\text{exp} = 000\dots 0, \text{frac} \neq 000\dots 0$
 - Numbers very close to 0.0
 - Lose precision as get smaller
 - Equispaced

Special Values

- Condition: $\text{exp} = 111\dots 1$
- Case: $\text{exp} = 111\dots 1, \text{frac} = 000\dots 0$
 - Represents value ∞ (infinity)
 - Operation that overflows
 - Both positive and negative
 - E.g., $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$
- Case: $\text{exp} = 111\dots 1, \text{frac} \neq 000\dots 0$
 - Not-a-Number (NaN)
 - Represents case when no numeric value can be determined
 - E.g., $\text{sqrt}(-1)$, $\infty - \infty$, $\infty \times 0$


Visualization: Floating Point Encodings



Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
 - Rounding, addition, multiplication
 - Floating point in C
 - Summary

Tiny Floating Point Example



► 8-bit Floating Point Representation

- the sign bit is in the most significant bit
- the next four bits are the exponent, with a bias of 7
- the last three bits are the *frac*

► Same general form as IEEE Format

- normalized, denormalized
- representation of 0, NaN, infinity


Dynamic Range (Positive Only)

	s	exp	frac	E	Value	
Denormalized numbers	0	0000	000	-6	0	closest to zero
	0	0000	001	-6	$1/8 * 1/64 = 1/512$	
	0	0000	010	-6	$2/8 * 1/64 = 2/512$	
	...					
	0	0000	110	-6	$6/8 * 1/64 = 6/512$	largest denorm
Normalized numbers	0	0000	111	-6	$7/8 * 1/64 = 7/512$	
	0	0001	000	-6	$8/8 * 1/64 = 8/512$	
	0	0001	001	-6	$9/8 * 1/64 = 9/512$	smallest norm
	...					
	0	0110	110	-1	$14/8 * 1/2 = 14/16$	closest to 1 below
Normalized numbers	0	0110	111	-1	$15/8 * 1/2 = 15/16$	
	0	0111	000	0	$8/8 * 1 = 1$	
	0	0111	001	0	$9/8 * 1 = 9/8$	closest to 1 above
	0	0111	010	0	$10/8 * 1 = 10/8$	
	...					
Normalized numbers	0	1110	110	7	$14/8 * 128 = 224$	largest norm
	0	1110	111	7	$15/8 * 128 = 240$	
	0	1111	000	n/a	inf	


Distribution of Values

► 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- Bias is $2^{3-1}-1 = 3$




► Notice how the distribution gets denser toward zero.



Distribution of Values (close-up view)

► 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- Bias is 3



Interesting Numbers

Description	exp	frac	Numeric Value
▶ Zero	00...00	00...00	0.0
▶ Smallest Pos. Denorm.	00...00	00...01	$2^{-(23,52)} \times 2^{-(126,1022)}$
◦ Single $\approx 1.4 \times 10^{-45}$			
◦ Double $\approx 4.9 \times 10^{-324}$			
▶ Largest Denormalized	00...00	11...11	$(1.0 - \epsilon) \times 2^{-(126,1022)}$
◦ Single $\approx 1.18 \times 10^{-38}$			
◦ Double $\approx 2.2 \times 10^{-308}$			
▶ Smallest Pos. Normalized	00...01	00...00	$1.0 \times 2^{-(126,1022)}$
◦ Just larger than largest denormalized			
▶ One	01...11	00...00	1.0
▶ Largest Normalized	11...10	11...11	$(2.0 - \epsilon) \times 2^{(127,1023)}$
◦ Single $\approx 3.4 \times 10^{38}$			
◦ Double $\approx 1.8 \times 10^{308}$			

Special Properties of Encoding

- ▶ FP Zero Same as Integer Zero
 - All bits = 0
- ▶ Can (Almost) Use Unsigned Integer Comparison
 - Must first compare sign bits
 - Must consider $-0 = 0$
 - NaNs problematic
 - Will be greater than any other values
 - What should comparison yield?
 - Otherwise OK
 - Denorm vs. normalized
 - Normalized vs. infinity

Today: Floating Point

- ▶ Background: Fractional binary numbers
- ▶ IEEE floating point standard: Definition
- ▶ Example and properties
- ▶ Rounding, addition, multiplication
- ▶ Floating point in C
- ▶ Summary

Floating Point Operations: Basic Idea

- ▶ $x +_f y = \text{Round}(x + y)$
- ▶ $x \times_f y = \text{Round}(x \times y)$
- ▶ Basic idea
 - First **compute exact result**
 - Make it fit into desired precision
 - Possibly overflow if exponent too large
 - Possibly **round to fit into** frac

Rounding

- ▶ Rounding Modes (illustrate with \$ rounding)

	\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
◦ Towards zero	\$1	\$1	\$1	\$2	-\$1
◦ Round down ($-\infty$)	\$1	\$1	\$1	\$2	-\$2
◦ Round up ($+\infty$)	\$2	\$2	\$2	\$3	-\$1
◦ Nearest Even (default)	\$1	\$2	\$2	\$2	-\$2

- ▶ What are the advantages of the modes?

Closer Look at Round-To-Even

- ▶ Default Rounding Mode

- Hard to get any other kind without dropping into assembly
- All others are statistically biased
 - Sum of set of positive numbers will consistently be over- or under-estimated

- ▶ Applying to Other Decimal Places / Bit Positions

- When exactly halfway between two possible values
 - Round so that least significant digit is even
- E.g., round to nearest hundredth

1.2349999	1.23	(Less than half way)
1.2350001	1.24	(Greater than half way)
1.2350000	1.24	(Half way—round up)
1.2450000	1.24	(Half way—round down)

Rounding Binary Numbers

- ▶ Binary Fractional Numbers

- “Even” when least significant bit is 0
- “Half way” when bits to right of rounding position = 100...2

- ▶ Examples

- Round to nearest 1/4 (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded
2 3/32	10.000 11 ₂	10.00 ₂	(<1/2—down)	2
2 3/16	10.001 10 ₂	10.01 ₂	(>1/2—up)	2 1/4
2 7/8	10.111 00 ₂	11.00 ₂	(1/2—up)	3
2 5/8	10.101 00 ₂	10.10 ₂	(1/2—down)	2 1/2

FP Multiplication

- ▶ $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$

- ▶ Exact Result: $(-1)^s M 2^E$

- Sign s: $s_1 \wedge s_2$
- Significand M: $M_1 \times M_2$
- Exponent E: $E_1 + E_2$

- ▶ Fixing

- If $M \geq 2$, shift M right, increment E
- If E out of range, overflow
- Round M to fit `frac` precision

- ▶ Implementation

- Biggest chore is multiplying significands


Floating Point Addition

▶ $(-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$

◦ Assume $E_1 > E_2$

▶ Exact Result: $(-1)^s M 2^E$

- Sign s , significand M :
 - Result of signed align & add
- Exponent E : E_1



▶ Fixing

- If $M \geq 2$, shift M right, increment E
- if $M < 1$, shift M left k positions, decrement E by k
- Overflow if E out of range
- Round M to fit $frac$ precision

Mathematical Properties of FP Add

▶ Compare to those of Abelian Group

- Closed under addition? **Yes**
 - But may generate infinity or NaN
- Commutative? **Yes**
- Associative? **No**
 - Overflow and inexactness of rounding
- 0 is additive identity? **Yes**
- Every element has additive inverse **Almost**
 - Except for infinities & NaNs

▶ Monotonicity

- $a \geq b \Rightarrow a+c \geq b+c$? **Almost**
 - Except for infinities & NaNs

Mathematical Properties of FP Mult

▶ Compare to Commutative Ring

- Closed under multiplication?
 - But may generate infinity or NaN
- Multiplication Commutative?
- Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
- 1 is multiplicative identity?
- Multiplication distributes over addition?
 - Possibility of overflow, inexactness of rounding

▶ Monotonicity

- $a \geq b$ & $c \geq 0 \Rightarrow a * c \geq b * c$?
 - Except for infinities & NaNs

Today: Floating Point

- ▶ Background: Fractional binary numbers
- ▶ IEEE floating point standard: Definition
- ▶ Example and properties
- ▶ Rounding, addition, multiplication
- ▶ Floating point in C
- ▶ Summary

Floating Point in C

- ▶ C Guarantees Two Levels
 - float single precision
 - double double precision
- ▶ Conversions/Casting
 - Casting between int, float, and double changes bit representation
 - double/float → int
 - Truncates fractional part
 - Like rounding toward zero
 - Not defined when out of range or NaN: Generally sets to TMin
 - int → double
 - Exact conversion, as long as int has ≤ 53 bit word size
 - int → float
 - Will round according to rounding mode

Floating Point Puzzles

- ▶ For each of the following C expressions, either:
 - Argue that it is true for all argument values
 - Explain why not true

```
int x = ...;
float f = ...;
double d = ...;
```

Assume neither
d nor f is NaN

- $x == (int)(float) x$
- $x == (int)(double) x$
- $f == (float)(double) f$
- $d == (float) d$
- $f == -(-f);$
- $2/3 == 2/3.0$
- $d < 0.0 \Rightarrow ((d*2) < 0.0)$
- $d > f \Rightarrow -f > -d$
- $d * d \geq 0.0$
- $(d+f)-d == f$

Summary

- ▶ IEEE Floating Point has clear mathematical properties
- ▶ Represents numbers of form $M \times 2^E$
- ▶ One can reason about operations independent of implementation
 - As if computed with perfect precision and then rounded
- ▶ Not the same as real arithmetic
 - Violates associativity/distributivity
 - Makes life difficult for compilers & serious numerical applications programmers

More Slides

Creating Floating Point Number

Steps


- Normalize to have leading 1
- Round to fit within fraction
- Postnormalize to deal with effects of rounding

Case Study


- Convert 8-bit unsigned numbers to tiny floating point format

Example Numbers

128	10000000
15	00001101
33	00010001
35	00010011
138	10001010
63	00111111



Normalize



Requirement

- Set binary point so that numbers of form 1.xxxxx
- Adjust all to have leading one
 - Decrement exponent as shift left

Value	Binary	Fraction	Exponent
128	10000000	1.0000000	7
15	00001101	1.1010000	3
17	00010001	1.0001000	4
19	00010011	1.0011000	4
138	10001010	1.0001010	7
63	00111111	1.1111100	5

Rounding

1.BBG**R**XXX

Guard bit: LSB of result

Round bit: 1st bit removed

Sticky bit: OR of remaining bits

Round up conditions

- Round = 1, Sticky = 1 \Rightarrow > 0.5
- Guard = 1, Round = 1, Sticky = 0 \Rightarrow Round to even

Value	Fraction	GRS	Incr?	Rounded
128	1.0000000	000	N	1.000
15	1.1010000	100	N	1.101
17	1.0001000	010	N	1.000
19	1.0011000	110	Y	1.010
138	1.0001010	011	Y	1.001
63	1.1111100	111	Y	10.000

Postnormalize

Issue

- Rounding may have caused overflow
- Handle by shifting right once & incrementing exponent

Value	Rounded	Exp	Adjusted	Result
128	1.000	7		128
15	1.101	3		15
17	1.000	4		16
19	1.010	4		20
138	1.001	7		134
63	10.000	5	1.000/6	64