

## Machine Level Programming IV Data and Arrays

CS 341: Intro. to Computer  
Architecture & Organization

Andree Jacobson

### Today

- Arrays
  - One-dimensional
  - Multi-dimensional (nested)
  - Multi-level
- Structures

### Basic Data Types

#### ► Integral

- Stored & operated on in general (integer) registers
- Signed vs. unsigned depends on instructions used

Intel	ASM	Bytes	C
byte	b	1	[unsigned] char
word	w	2	[unsigned] short
double word	l	4	[unsigned] int
quad word	q	8	[unsigned] long int (x86-64)

#### ► Floating Point

- Stored & operated on in floating point registers

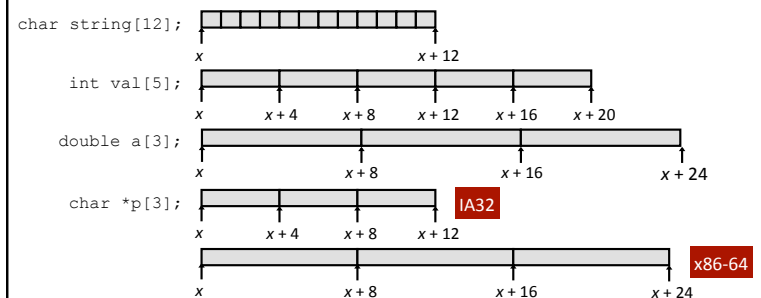
Intel	ASM	Bytes	C
Single	s	4	float
Double	l	8	double
Extended	t	10/12/16	long double

### Array Allocation

#### ► Basic Principle

$T \ A[L];$

- Array of data type  $T$  and length  $L$
- Contiguously allocated region of  $L * \text{sizeof}(T)$  bytes

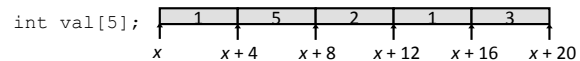


## Array Access

### Basic Principle

$T \ A[L];$

- Array of data type  $T$  and length  $L$
- Identifier  $A$  can be used as a pointer to array element 0: Type  $T^*$



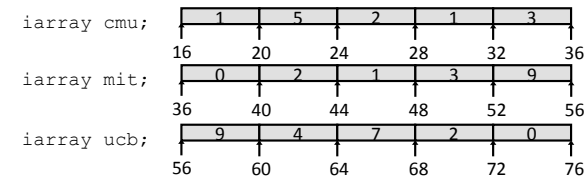
### Reference Type Value

`val[4]`  
`val`  
`val+1`  
`&val[2]`  
`val[5]`  
`*(val+1)`  
`val + i`

## Array Example

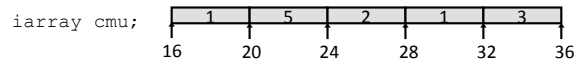
```
#define ALEN 5
typedef int iarray[ZLEN];

iarray cmu = { 1, 5, 2, 1, 3 };
iarray mit = { 0, 2, 1, 3, 9 };
iarray ucb = { 9, 4, 7, 2, 0 };
```



- Declaration “`iarray cmu`” equivalent to “`int cmu[5]`”
- Example arrays were allocated in successive 20 byte blocks
  - Not guaranteed to happen in general

## Array Accessing Example



```
int get_digit
(iarray z, int dig)
{
    return z[dig];
}
```

IA32

```
# %edx = z
# %eax = dig
movl (%edx,%eax,4),%eax # z[dig]
```

- Register `%edx` contains starting address of array
- Register `%eax` contains array index
- Desired digit at  $4 * \text{array index} + \text{starting address}$
- Use memory reference  $(\text{starting address}, \text{array index}, 4)$

## Array Loop Example (IA32)

```
void incr(iarray a) {
    int i;
    for (i = 0; i < ALEN; i++)
        a[i]++;
}
```

```
# edx = a
movl $0, %eax # %eax = i
.L4: # loop:
    addl $1, (%edx,%eax,4) # a[i]++
    addl $1, %eax # i++
    cmpl $5, %eax # i:5
    jne .L4 # if !=, goto loop
```

## Pointer Loop Example (IA32)

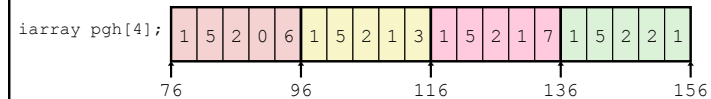
```
void incr_p(iarray a) {
    int *aend = a+ALEN;
    do {
        (*a)++;
        a++;
    } while (a != aend);
}
```

```
void incr_v(iarray a) {
    void *va = a;
    int i = 0;
    do {
        (*(int *) (va+i))++;
        i += ISIZE;
    } while (i != ISIZE*ALEN);
}
```

```
# edx = a = va
movl $0, %eax          # i = 0
.L8:                   # loop:
    addl $1, (%edx,%eax) # Increment va+i
    addl $4, %eax        # i += 4
    cmpl $20, %eax       # Compare i:20
    jne .L8              # if !=, goto loop
```

## Nested Array Example

```
#define PCOUNT 4
iarray pgh[PCOUNT] =
    {{1, 5, 2, 0, 6},
     {1, 5, 2, 1, 3},
     {1, 5, 2, 1, 7},
     {1, 5, 2, 2, 1}};
```



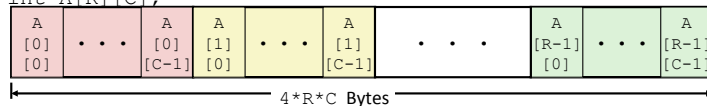
- ▶ “iarray pgh[4]” equivalent to “int pgh[4][5]”
  - Variable **pgh**: array of 4 elements, allocated contiguously
  - Each element is an array of 5 **int**'s, allocated contiguously
- ▶ “Row-Major” ordering of all elements guaranteed

## Multidimensional (Nested) Arrays

- ▶ Declaration
  - $T \ A[R][C];$
  - 2D array of type  $T$
  - $R$  rows,  $C$  columns
  - $\text{sizeof}(\text{Type } T): K$  bytes
- ▶ Array Size
  - $R * C * K$  bytes
- ▶ Arrangement
  - Row-Major Ordering

$A[0][0]$	...	$A[0][C-1]$
$\vdots$		$\vdots$
$A[R-1][0]$	...	$A[R-1][C-1]$

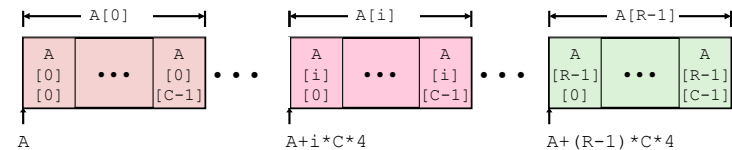
```
int A[R][C];
```



## Nested Array Row Access

- ▶ Row Vectors
  - $A[i]$  is array of  $C$  elements
  - Each element of type  $T$  requires  $K$  bytes
  - Starting address  $A + i * (C * K)$

```
int A[R][C];
```



## Nested Array Row Access Code

```
int *get_pgh(int index)
{
    return pgh[index];
}
```

```
#define PCOUNT 4
iarray pgh[PCOUNT] =
{ {1, 5, 2, 0, 6},
  {1, 5, 2, 1, 3 },
  {1, 5, 2, 1, 7 },
  {1, 5, 2, 2, 1 } };
```

```
# %eax = index
leal (%eax,%eax,4),%eax # 5 * index
leal pgh(,%eax,4),%eax # pgh + (20 * index)
```

### ► Row Vector

- `pgh[index]` is array of 5 int's
- Starting address `pgh+20*index`

### ► IA32 Code

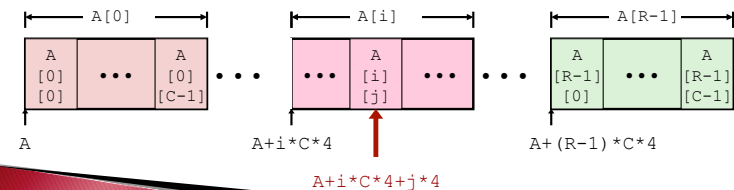
- Computes and returns address
- Compute as `pgh + 4*(index+4*index)`

## Nested Array Element Access

### ► Array Elements

- $A[i][j]$  is element of type  $T$ , which requires  $K$  bytes
- Address  $A + i * (C * K) + j * K = A + (i * C + j) * K$

```
int A[R][C];
```



## Nested Array Element Access Code

```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

```
movl 8(%ebp), %eax # index
leal (%eax,%eax,4), %eax # 5*index
addl 12(%ebp), %eax # 5*index+dig
movl pgh(,%eax,4), %eax # offset 4*(5*index+dig)
```

### ► Array Elements

- `pgh[index][dig]` is int
- Address: `pgh + 20*index + 4*dig`  
 $= pgh + 4*(5*index + dig)$

### ► IA32 Code

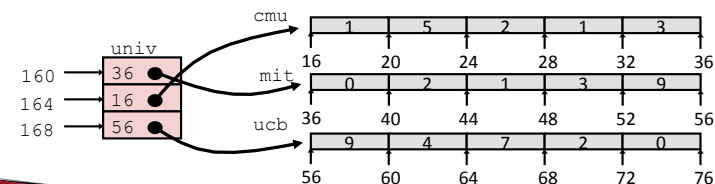
- Computes address `pgh + 4*((index+4*index)+dig)`

## Multi-Level Array Example

```
iarray cmu = { 1, 5, 2, 1, 3 };
iarray mit = { 0, 2, 1, 3, 9 };
iarray ucb = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3
int *univ[UCOUNT] = {mit, cmu, ucb};
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer
  - 4 bytes
- Each pointer points to array of int's



## Element Access in Multi-Level Array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```

```
movl 8(%ebp), %eax    # index
movl univ(,%eax,4), %edx # p = univ[index]
movl 12(%ebp), %eax    # dig
movl (%edx,%eax,4), %eax # p[dig]
```

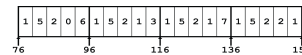
### ► Computation (IA32)

- Element access **Mem[Mem[univ+4\*index]+4\*dig]**
- Must do two memory reads
  - First get pointer to row array
  - Then access element within array

## Array Element Accesses

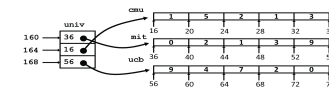
### Nested array

```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```



### Multi-level array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```



Accesses look similar in C, but addresses very different:

Mem[pgh+20\*index+4\*dig]

Mem[Mem[univ+4\*index]+4\*dig]

## N X N Matrix Code

### ► Fixed dimensions

- Know value of N at compile time

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele
(fix_matrix a, int i, int j)
{
    return a[i][j];
}
```

### ► Variable dimensions, explicit indexing

- Traditional way to implement dynamic arrays

```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element a[i][j] */
int vec_ele
(int n, int *a, int i, int j)
{
    return a[IDX(n,i,j)];
}
```

### ► Variable dimensions, implicit indexing

- Now supported by gcc

```
/* Get element a[i][j] */
int var_ele
(int n, int a[n][n], int i, int j) {
    return a[i][j];
}
```

## 16 X 16 Matrix Access

### ■ Array Elements

- Address  $A + i * (C * K) + j * K$
- $C = 16, K = 4$

```
/* Get element a[i][j] */
int fix_ele(fix_matrix a, int i, int j) {
    return a[i][j];
}
```

```
movl 12(%ebp), %edx    # i
sall $6, %edx          # i*64
movl 16(%ebp), %eax    # j
sall $2, %eax          # j*4
addl 8(%ebp), %eax     # a + j*4
movl (%eax,%edx), %eax # *(a + j*4 + i*64)
```

## n X n Matrix Access

### ■ Array Elements

- Address  $A + i * (C * K) + j * K$
- $C = n, K = 4$

```
/* Get element a[i][j] */
int var_ele(int n, int a[n][n], int i, int j) {
    return a[i][j];
}
```

```
movl 8(%ebp), %eax    # n
sall $2, %eax         # n*4
movl %eax, %edx       # n*4
imull 16(%ebp), %edx  # i*n*4
movl 20(%ebp), %eax   # j
sall $2, %eax         # j*4
addl 12(%ebp), %eax   # a + j*4
movl (%eax,%edx), %eax # *(a + j*4 + i*n*4)
```

## Optimizing Fixed Array Access



```
#define N 16
typedef int fix_matrix[N][N];
```

- Computation
  - Step through all elements in column j
- Optimization
  - Retrieving successive elements from single column

```
/* Retrieve column j from array */
void fix_column
(fix_matrix a, int j, int *dest)
{
    int i;
    for (i = 0; i < N; i++)
        dest[i] = a[i][j];
}
```

## Optimizing Fixed Array Access

### ► Optimization

- Compute  $ajp = \&a[i][j]$ 
  - Initially =  $a + 4*j$
  - Increment by  $4*N$

Register	Value
%ecx	ajp
%ebx	dest
%edx	i

```
/* Retrieve column j from array */
void fix_column
(fix_matrix a, int j, int *dest)
{
    int i;
    for (i = 0; i < N; i++)
        dest[i] = a[i][j];
}
```

```
.L8:      # loop:
movl    (%ecx), %eax    # Read *ajp
movl    %eax, (%ebx,%edx,4) # Save in dest[i]
addl    $1, %edx        # i++
addl    $64, %ecx       # ajp += 4*N
cmpl    $16, %edx       # i:N
jne     .L8             # if !=, goto loop
```

## Optimizing Variable Array Access

- Compute  $ajp = \&a[i][j]$ 
  - Initially =  $a + 4*j$
  - Increment by  $4*n$

Register	Value
%ecx	ajp
%edi	dest
%edx	i
%ebx	4*n
%esi	n

```
/* Retrieve column j from array */
void var_column
(int n, int a[n][n],
 int j, int *dest)
{
    int i;
    for (i = 0; i < n; i++)
        dest[i] = a[i][j];
}
```

```
.L18:      # loop:
movl    (%ecx), %eax    # Read *ajp
movl    %eax, (%edi,%edx,4) # Save in dest[i]
addl    $1, %edx        # i++
addl    %ebx, %ecx       # ajp += 4*n
cmpl    %edx, %esi       # n:i
jg      .L18             # if >, goto loop
```

## Today

### ■ Arrays

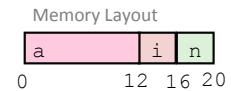
- One-dimensional
- Multi-dimensional (nested)
- Multi-level

### ■ Structures

- Allocation
- Access

## Structure Allocation

```
struct rec {
    int a[3];
    int i;
    struct rec *n;
};
```

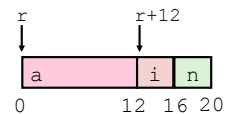


### ► Concept

- Contiguously-allocated region of memory
- Refer to members within structure by names
- Members may be of different types

## Structure Access

```
struct rec {
    int a[3];
    int i;
    struct rec *n;
};
```



### ► Accessing Structure Member

- Pointer indicates first byte of structure
- Access elements with offsets

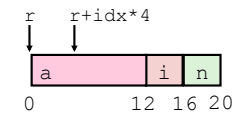
```
void
set_i(struct rec *r,
      int val)
{
    r->i = val;
}
```

IA32 Assembly

```
# %edx = val
# %eax = r
movl %edx, 12(%eax) # Mem[r+12] = val
```

## Generating Pointer to Structure Member

```
struct rec {
    int a[3];
    int i;
    struct rec *n;
};
```



### ► Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Arguments
  - Mem[%ebp+8]: **r**
  - Mem[%ebp+12]: **idx**

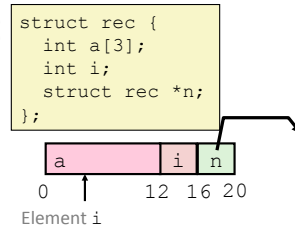
```
int *get_ap
(struct rec *r, int idx)
{
    return &r->a[idx];
}
```

```
movl 12(%ebp), %eax # Get idx
sall $2, %eax       # idx*4
addl 8(%ebp), %eax  # r+idx*4
```

## Following Linked List

### ► C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->n;
    }
}
```

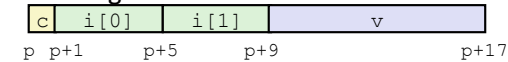


Register	Value
%edx	r
%ecx	val

```
.L17:      movl    12(%edx), %eax      # loop:
          # r->i
          movl    %ecx, (%edx,%eax,4) # r->a[i] = val
          movl    16(%edx), %edx      # r = r->n
          testl   %edx, %edx          # Test r
          jne     .L17                # If != 0 goto loop
```

## Structures & Alignment

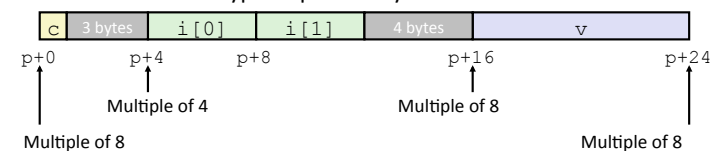
### ► Unaligned Data



```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

### ► Aligned Data

- Primitive data type requires K bytes



## Alignment Principles

### ► Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K
- Required on some machines; advised on IA32
  - treated differently by IA32 Linux, x86-64 Linux, and Windows!

### ► Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
  - Inefficient to load or store datum that spans quad word boundaries
  - Virtual memory very tricky when datum spans 2 pages

### ► Compiler

- Inserts gaps in structure to ensure correct alignment of fields

## Specific Cases of Alignment (IA32)

- 1 byte: char, ...
  - no restrictions on address
- 2 bytes: short, ...
  - lowest 1 bit of address must be 0<sub>2</sub>
- 4 bytes: int, float, char \*, ...
  - lowest 2 bits of address must be 00<sub>2</sub>
- 8 bytes: double, ...
  - Windows (and most other OS's & instruction sets):
    - lowest 3 bits of address must be 000<sub>2</sub>
  - Linux:
    - lowest 2 bits of address must be 00<sub>2</sub>
    - i.e., treated the same as a 4-byte primitive data type
- 12 bytes: long double
  - Windows, Linux:
    - lowest 2 bits of address must be 00<sub>2</sub>
    - i.e., treated the same as a 4-byte primitive data type



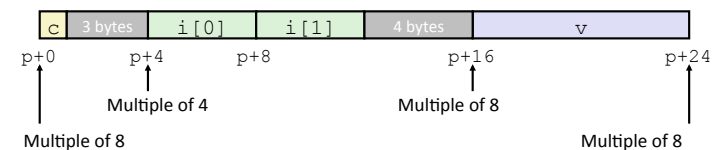
## Specific Cases of Alignment (x86-64)

- ▶ 1 byte: char, ...
  - no restrictions on address
- ▶ 2 bytes: short, ...
  - lowest 1 bit of address must be 0<sub>2</sub>
- ▶ 4 bytes: int, float, ...
  - lowest 2 bits of address must be 00<sub>2</sub>
- ▶ 8 bytes: double, char \*, ...
  - Windows & Linux:
    - lowest 3 bits of address must be 000<sub>2</sub>
- ▶ 16 bytes: long double
  - Linux:
    - lowest 3 bits of address must be 000<sub>2</sub>
    - i.e., treated the same as a 8-byte primitive data type

## Satisfying Alignment with Structures

- ▶ Within structure:
  - Must satisfy each element's alignment requirement
- ▶ Overall structure placement
  - Each structure has alignment requirement K
    - K = Largest alignment of any element
  - Initial address & structure length must be multiples of K
- ▶ Example (under Windows or x86-64):
  - K = 8, due to double element

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```



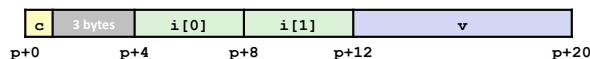
## Different Alignment Conventions

- ▶ x86-64 or IA32 Windows:
  - K = 8, due to double element

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```



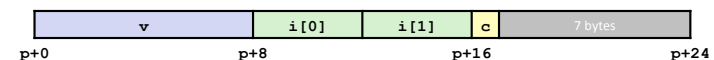
- ▶ IA32 Linux
  - K = 4; double treated like a 4-byte data type



## Meeting Overall Alignment Requirement

- ▶ For largest alignment requirement K
- ▶ Overall structure must be multiple of K

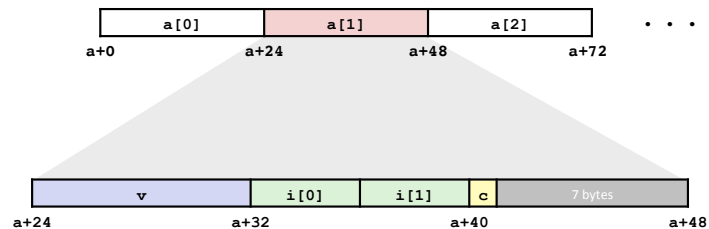
```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```



## Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element

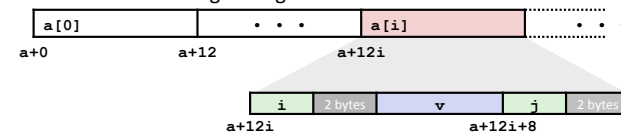
```
struct S2 {
    double v;
    int i[2];
    char c;
} a[10];
```



## Accessing Array Elements

- Compute array offset  $12i$ 
  - $\text{sizeof}(S3)$ , including alignment spacers
- Element  $j$  is at offset 8 within structure
  - Assembler gives offset  $a+8$ 
    - Resolved during linking

```
struct S3 {
    short i;
    float v;
    short j;
} a[10];
```



```
short get_j(int idx)
{
    return a[idx].j;
}
```

```
# %eax = idx
leal (%eax,%eax,2),%eax # 3*idx
movswl a+8(,%eax,4),%eax
```

## Saving Space

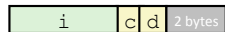
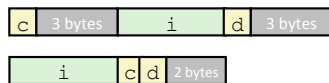
- Put large data types first

```
struct S4 {
    char c;
    int i;
    char d;
} *p;
```



```
struct S5 {
    int i;
    char c;
    char d;
} *p;
```

- Effect ( $K=4$ )

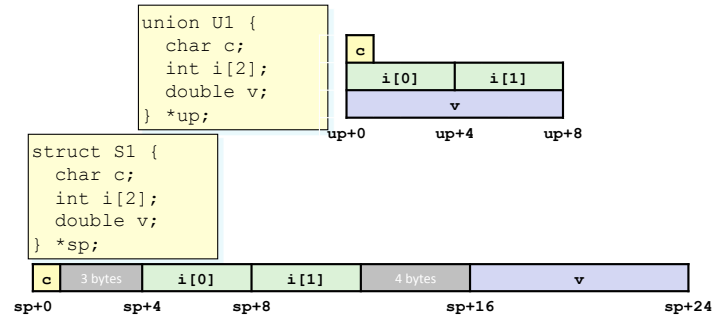


## Today

- Structures
  - Alignment
- Unions
- Memory Layout
- Buffer Overflow
  - Vulnerability
  - Protection

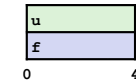
## Union Allocation

- Allocate according to largest element
- Should (probably) only use one field at a time!



## Using Union to Access Bit Patterns

```
typedef union {
    float f;
    unsigned u;
} bit_float_t;
```



```
float bit2float(unsigned u)
{
    bit_float_t arg;
    arg.u = u;
    return arg.f;
}
```

Same as (float) u?

```
unsigned float2bit(float f)
{
    bit_float_t arg;
    arg.f = f;
    return arg.u;
}
```

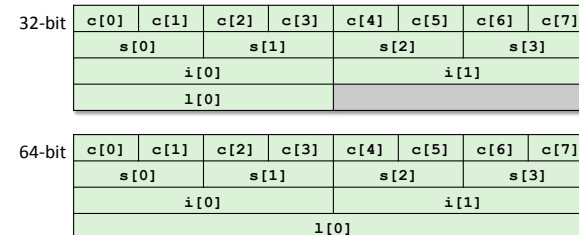
Same as (unsigned) f?

## Byte Ordering Revisited

- Idea
  - Short/long/quad words stored in memory as 2/4/8 consecutive bytes
  - Which is most (least) significant?
  - Can cause problems when exchanging binary data between machines
- Big Endian
  - Most significant byte has lowest address
  - Sparc
- Little Endian
  - Least significant byte has lowest address
  - Intel x86

## Byte Ordering Example

```
union {
    unsigned char c[8];
    unsigned short s[4];
    unsigned int i[2];
    unsigned long l[1];
} dw;
```



## Byte Ordering Example (Cont).

```
int j;
for (j = 0; j < 8; j++)
    dw.c[j] = 0xf0 + j;

printf("Characters 0-7 == [0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x]\n",
       dw.c[0], dw.c[1], dw.c[2], dw.c[3], dw.c[4], dw.c[5], dw.c[6], dw.c[7]);

printf("Shorts 0-3 == [0x%x,0x%x,0x%x,0x%x]\n",
       dw.s[0], dw.s[1], dw.s[2], dw.s[3]);

printf("Ints 0-1 == [0x%x,0x%x]\n",
       dw.i[0], dw.i[1]);

printf("Long 0 == [0x%x]\n",
       dw.l[0]);
```

## Byte Ordering on IA32

Little Endian

f0	f1	f2	f3	f4	f5	f6	f7
c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

LSB

MSB

LSB

MSB

Print

Output:

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]  
Shorts 0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]  
Ints 0-1 == [0xf3f2f1f0,0xf7f6f5f4]  
Long 0 == [0xf3f2f1f0]

## Byte Ordering on Sun

Big Endian

f0	f1	f2	f3	f4	f5	f6	f7
c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

MSB → Print → LSB    MSB → LSB

Output on Sun:

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]  
Shorts 0-3 == [0xf0f1,0xf2f3,0xf4f5,0xf6f7]  
Ints 0-1 == [0xf0f1f2f3,0xf4f5f6f7]  
Long 0 == [0xf0f1f2f3]

## Byte Ordering on x86-64

Little Endian

f0	f1	f2	f3	f4	f5	f6	f7
c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

LSB ← Print → MSB

Output on x86-64:

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]  
Shorts 0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]  
Ints 0-1 == [0xf3f2f1f0,0xf7f6f5f4]  
Long 0 == [0xf7f6f5f4f3f2f1f0]