

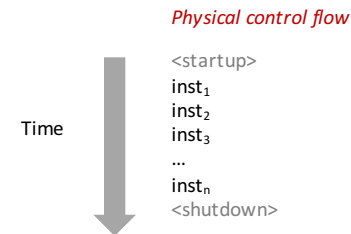
Exceptional Control Flow

CS 341: Intro. to Computer
Architecture & Organization

Prof. Andree Jacobson

Control Flow

- ▶ Processors do only one thing:
 - From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
 - This sequence is the CPU's *control flow* (or *flow of control*)



Altering the Control Flow


- ▶ Up to now: two mechanisms for changing control flow:
 - Jumps and branches
 - Call and return
 Both react to changes in *program state*
- ▶ Insufficient for a useful system:
Difficult to react to changes in *system state*
 - data arrives from a disk or a network adapter
 - instruction divides by zero
 - user hits Ctrl-C at the keyboard
 - System timer expires
- ▶ System needs mechanisms for “exceptional control flow”

Exceptional Control Flow

- ▶ Exists at all levels of a computer system
- ▶ Low level mechanisms
 - Exceptions
 - change in control flow in response to a system event (i.e., change in system state)
 - Combination of hardware and OS software
- ▶ Higher level mechanisms
 - Process context switch
 - Signals
 - Nonlocal jumps: setjmp()/longjmp()
 - Implemented by either:
 - OS software (context switch and signals)
 - C language runtime library (nonlocal jumps)

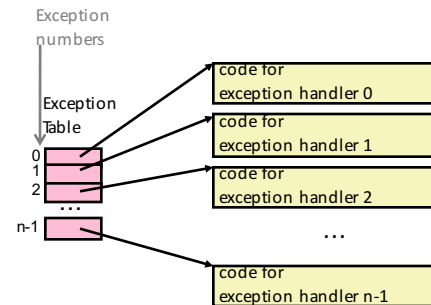
Exceptions

- ▶ An **exception** is a transfer of control to the OS in response to some **event** (i.e., change in processor state)



- ▶ Examples:
div by 0, arithmetic overflow, page fault, I/O request completes, Ctrl-C

Interrupt Vectors



- ▶ Each type of event has a unique exception number k
- ▶ k = index into exception table (a.k.a. interrupt vector)
- ▶ Handler k is called each time exception k occurs

Asynchronous Exceptions (Interrupts)

- ▶ Caused by events external to the processor
 - Indicated by setting the processor's interrupt pin
 - Handler returns to "next" instruction
- ▶ Examples:
 - I/O interrupts
 - hitting Ctrl-C at the keyboard
 - arrival of a packet from a network
 - arrival of data from a disk
 - Hard reset interrupt
 - hitting the reset button
 - Soft reset interrupt
 - hitting Ctrl-Alt-Delete on a PC


Synchronous Exceptions

- ▶ Caused by events that occur from instruction execution:
 - **Traps**
 - Intentional
 - Examples: **system calls**, breakpoint traps, special instructions
 - Returns control to "next" instruction
 - **Faults**
 - Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
 - Either re-executes faulting ("current") instruction or aborts
 - **Aborts**
 - Unintentional and unrecoverable
 - Examples: parity error, machine check
 - Aborts current program

Trap Example: Opening File

- ▶ User calls: `open(filename, options)`
- ▶ Function `open` executes system call instruction `int`
 - OS must find or create file, get it ready for reading or writing
 - Returns integer file descriptor

```
0804d070 <__libc_open>:
. . .
804d082: cd 80          int    $0x80
804d084: 5b            pop    %ebx
. . .
```

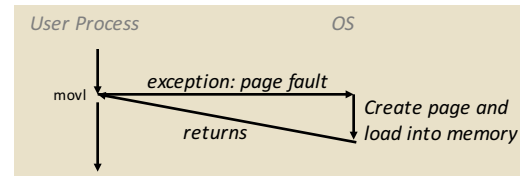


Fault Example: Page Fault

- ▶ User writes to memory location
- ▶ That portion (page) of user's memory is currently on disk

```
int a[1000];
main ()
{
    a[500] = 13;
}
```

```
80483b7: c7 05 10 9d 04 08 0d movl $0xd,0x8049d10
```

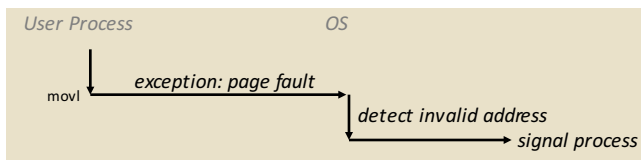


- ▶ Page handler must load page into physical memory
- ▶ Returns to faulting instruction
- ▶ Successful on second try

Fault Example: Invalid Memory Reference

```
int a[1000];
main ()
{
    a[5000] = 13;
}
```

```
80483b7: c7 05 60 e3 04 08 0d movl $0xd,0x804e360
```



- ▶ Page handler detects invalid address
- ▶ Sends SIGSEGV signal to user process
- ▶ User process exits with "segmentation fault"

Exception Table IA32 (Excerpt)


Exception Number	Description	Exception Class
0	Divide error	Fault
13	General protection fault	Fault
14	Page fault	Fault
18	Machine check	Abort
32-127	OS-defined	Interrupt or trap
128 (0x80)	System call	Trap
129-255	OS-defined	Interrupt or trap

Check Table 6-1:

<http://download.intel.com/design/processor/manuals/253665.pdf>

System Calls: Requesting OS Service

- OS uses dual modes to protect itself
 - User mode for unprivileged instructions
 - Kernel mode for privileged instructions
 - Mode bit** provided by hardware



Application View of OS

Applications can access OS via system calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

14

Today


- Exceptional Control Flow
- Processes

Processes

- Definition: A *process* is an instance of a running program.
 - One of the most profound ideas in computer science
 - Not the same as "program" or "processor"
- Process provides each program with two key abstractions:
 - Logical control flow
 - Each program seems to have exclusive use of the CPU
 - Private virtual address space
 - Each program seems to have exclusive use of main memory
- How are these illusions maintained?
 - Process executions interleaved (multitasking) or run on separate cores
 - Address spaces managed by virtual memory system


Concurrent Processes

- Two processes *run concurrently* (are concurrent) if their flows overlap in time
- Otherwise, they are *sequential*
- Examples (running on single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C



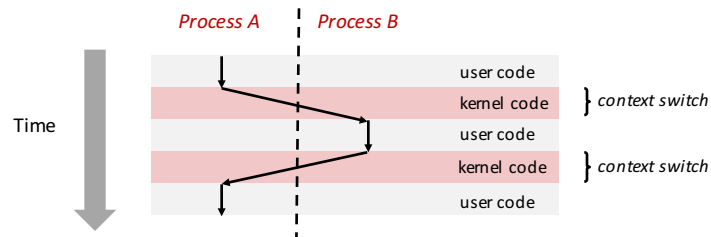
User View of Concurrent Processes

- Control flows for concurrent processes are physically disjoint in time
- However, we can think of concurrent processes as running in parallel with each other

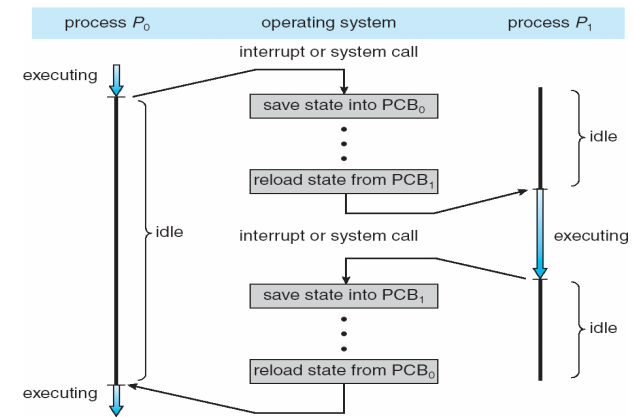


Context Switching

- Processes are managed by a shared chunk of OS code called the *kernel*
 - Important: the kernel is not a separate process, but rather runs as part of some user process
- Control flow passes from one process to another via a *context switch*

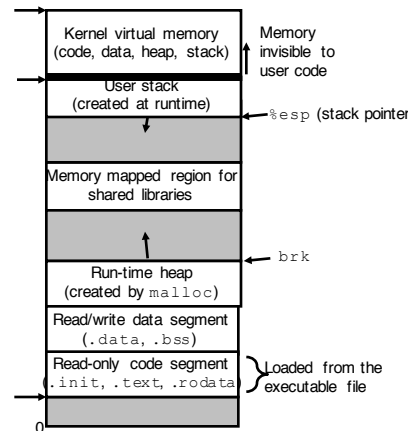


Context Switching (more details)



Address Spaces

- Processes represent a logical flow AND
- Processes represent a memory protection domain
 - aka address spaces
- Each process has its own private address space



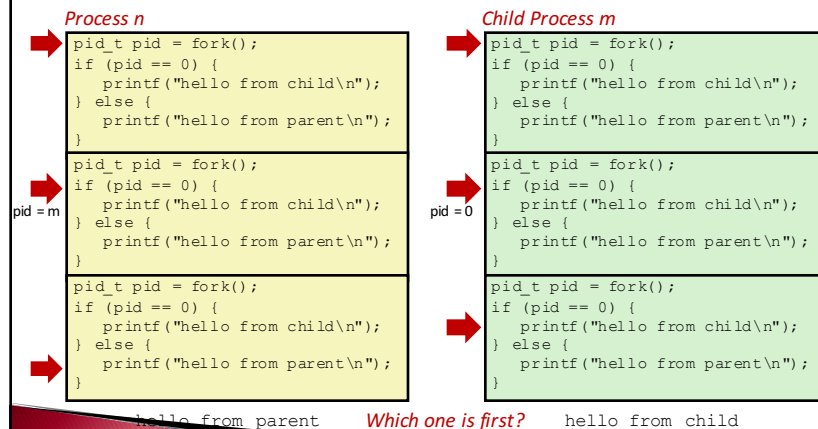
fork: Creating New Processes

- `int fork(void)`
 - creates a new process (child process) that is identical to the calling process (parent process)
 - returns 0 to the child process
 - returns child's `pid` to the parent process

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

- Fork is interesting (and often confusing) because it is called *once* but returns *twice*

Understanding fork



Fork Example #1


- Parent and child both run same code
 - Distinguish parent from child by return value from `fork`
- Start with same state, but each has private copy
 - Including shared output file descriptor
 - Relative ordering of their print statements undefined

```
void fork1 ()
{
    int x = 1;
    pid_t pid = fork();
    if (pid == 0) {
        printf("Child has x = %d\n", ++x);
    } else {
        printf("Parent has x = %d\n", --x);
    }
    printf("Bye from process %d with x = %d\n", getpid(), x);
}
```

Fork Example #2

- Both parent and child can continue forking


```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```



Fork Example #3

- Both parent and child can continue forking


```
void fork3()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("L2\n");
    fork();
    printf("Bye\n");
}
```



Fork Example #4

- Both parent and child can continue forking


```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
            fork();
        }
    }
    printf("Bye\n");
}
```



Fork Example #5

- Both parent and child can continue forking

```
void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
            fork();
        }
    }
    printf("Bye\n");
}
```



exit: Ending a process

- ▶ `void exit(int status)`
 - exits a process
 - Normally return with status 0
 - `atexit()` registers functions to be executed upon exit

```
void cleanup(void) {
    printf("cleaning up\n");
}

void fork6() {
    atexit(cleanup);
    fork();
    exit(0);
}
```

Zombies

- ▶ Idea
 - When process terminates, still consumes system resources
 - Various tables maintained by OS
 - Called a “zombie”
 - Living corpse, half alive and half dead
- ▶ Reaping
 - Performed by parent on terminated child
 - Parent is given exit status information
 - Kernel discards process
- ▶ What if parent doesn't reap?
 - If any parent terminates without reaping a child, then child will be reaped by `init` process
 - So, only need explicit reaping in long-running processes
 - e.g., shells and servers

Zombie Example

```
linux> ./forks 7 &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
linux> ps
  PID TTY          TIME CMD
 6585 ttty9      00:00:00 tcsh
 6639 ttty9      00:00:03 forks
 6640 ttty9      00:00:00 forks <defunct>
 6641 ttty9      00:00:00 ps
linux> kill 6639
[1] Terminated
linux> ps
  PID TTY          TIME CMD
 6585 ttty9      00:00:00 tcsh
 6642 ttty9      00:00:00 ps
```

```
void fork7()
{
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n",
            getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n",
            getpid());
        while (1)
            ; /* Infinite loop */
    }
}
```

- ▶ `ps` shows child process as “defunct”
- ▶ Killing parent allows child to be reaped by `init`

Nonterminating Child Example

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 ttty9      00:00:00 tcsh
 6676 ttty9      00:00:06 forks
 6677 ttty9      00:00:00 ps
linux> kill 6676
linux> ps
  PID TTY          TIME CMD
 6585 ttty9      00:00:00 tcsh
 6678 ttty9      00:00:00 ps
```

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n",
            getpid());
        while (1)
            ; /* Infinite loop */
    } else {
        printf("Terminating Parent, PID = %d\n",
            getpid());
        exit(0);
    }
}
```

- ▶ Child process still active even though parent has terminated
- ▶ Must kill explicitly, or else will keep running indefinitely


wait: Synchronizing with Children

- ▶ `int wait(int *child_status)`
 - suspends current process until one of its children terminates
 - return value is the **pid** of the child process that terminated
 - if **child_status != NULL**, then the object it points to will be set to a status indicating why the child process terminated

wait: Synchronizing with Children

```
void fork9() {
    int child_status;

    if (fork() == 0) {
        printf("HC: hello from child\n");
    }
    else {
        printf("HP: hello from parent\n");
        wait(&child_status);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
    exit();
}
```



wait() Example

- ▶ If multiple children completed, will take in arbitrary order
- ▶ Can use macros `WIFEXITED` and `WEXITSTATUS` to get information about exit status

```
void fork10()
{
    pid_t pid[N];
    int i;
    int child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = 0; i < N; i++) {
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}
```


waitpid(): Waiting for a Specific Process

- ▶ `waitpid(pid, &status, options)`
 - suspends current process until specific process terminates
 - various options (see textbook)

```
void fork11()
{
    pid_t pid[N];
    int i;
    int child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = N-1; i >= 0; i--) {
        pid_t wpid = waitpid(pid[i], &child_status, 0);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```


execve : Loading and Running Programs

- ```
int execve(
 char *filename,
 char *argv[],
 char *envp[]
)
```
- Loads and runs in current process:
    - Executable **filename**
    - With argument list **argv**
    - And environment variable list **envp**
  - Does not return (unless error)
  - Overwrites code, data, and stack
    - keeps pid, open files and signal context
  - Environment variables:
    - "name=value" strings
    - getenv and putenv



## execve Example

```
if ((pid = Fork()) == 0) { /* Child runs user job */
 if (execve(argv[0], argv, environ) < 0) {
 printf("%s: Command not found.\n", argv[0]);
 exit(0);
 }
}
```



## ECF Exists at All Levels of a System

- Exceptions
  - Hardware and operating system kernel software
- Process Context Switch
  - Hardware timer and kernel software
- Signals
  - Kernel software
- Nonlocal jumps
  - Application code

## Today

- Multitasking, shells
- Signals
- Nonlocal jumps

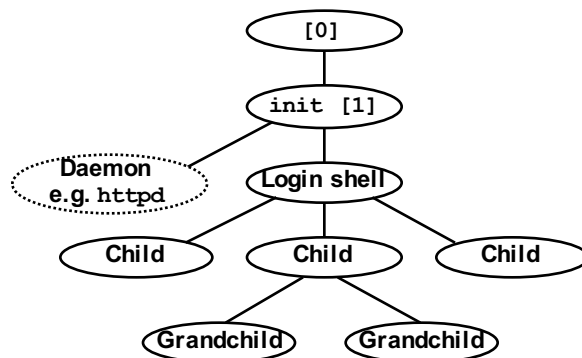
## The World of Multitasking

- ▶ System runs many processes concurrently
- ▶ Process: executing program
  - State includes memory image + register values (including program counter)
- ▶ Regularly switches from one process to another
  - Suspend process when it needs I/O resource or timer event occurs
  - Resume process when I/O available or given scheduling priority
- ▶ Appears to user(s) as if all processes executing simultaneously
  - Even though most systems can only execute one process at a time
  - Except possibly with lower performance than if running alone

## Programmer's Model of Multitasking

- ▶ Basic functions
  - **fork** spawns new process
    - Called once, returns twice
  - **exit** terminates own process
    - Called once, never returns
    - Puts it into “zombie” status
  - **wait** and **waitpid** wait for and reap terminated children
  - **execve** runs new program in existing process
    - Called once, (normally) never returns

## Unix Process Hierarchy



## What Is a “Background Job”?

- ▶ Users generally run one command at a time
  - Type command, read output, type another command
- ▶ Some programs run “for a long time”
 

```
unix> sleep 7200; rm /tmp/junk # shell stuck for 2 hours
```
- ▶ A “background” job is a process we don't want to wait for

```
unix> (sleep 7200 ; rm /tmp/junk) &
[1] 907
unix> # ready for next command
```

## Shell Programs

- ▶ A *shell* is an application program that runs programs on behalf of the user.

- **sh** Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
- **csh** BSD Unix C shell (**tcsh**: enhanced **csh** at CMU and elsewhere)
- **bash** "Bourne-Again" Shell

```
int main() {
 char cmdline[MAXLINE];

 while (1) {
 /* read */
 printf("> ");
 fgets(cmdline, MAXLINE, stdin);
 if (feof(stdin))
 exit(0);

 /* evaluate */
 eval(cmdline);
 }
}
```

*Execution is a sequence of read/evaluate steps*

## Simple Shell eval Function

```
void eval(char *cmdline) {
 char *argv[MAXARGS]; /* argv for execve() */
 int bg; /* should the job run in bg or fg? */
 pid_t pid; /* process id */

 bg = parseline(cmdline, argv);
 if (!builtin_command(argv)) {
 if ((pid = Fork()) == 0) { /* child runs user job */
 if (execve(argv[0], argv, environ) < 0) {
 printf("%s: Command not found.\n", argv[0]);
 exit(0);
 }
 }

 if (!bg) { /* parent waits for fg job to terminate */
 int status;
 if (waitpid(pid, &status, 0) < 0)
 unix_error("waitfg: waitpid error");
 }
 else /* otherwise, don't wait for bg job */
 printf("%d %s", pid, cmdline);
 }
}
```

## Problem with Simple Shell Example

- ▶ Our example shell correctly waits for and reaps foreground jobs
- ▶ But what about background jobs?
  - Will become zombies when they terminate
  - Will never be reaped because shell (typically) will not terminate
  - Will create a memory leak that could run the kernel out of memory
  - Modern Unix: once you exceed your process quota, your shell can't run any new commands for you: fork() returns -1

```
unix> limit maxproc # csh syntax
maxproc 202752
unix> ulimit -u # bash syntax
202752
```

## ECF to the Rescue!

- ▶ Problem
  - The shell doesn't know when a background job will finish
    - By nature, it could happen at any time
  - The shell's regular control flow can't reap exited background processes in a timely fashion
    - Regular control flow is "wait until running job completes, then reap it"
- ▶ Solution: Exceptional control flow
  - The kernel will interrupt regular processing to alert us when a child (background) process completes
  - In Unix, the alert mechanism is called a **signal**

## Signals

- ▶ A **signal** is a small message that notifies a process that an event of some type has occurred in the system
  - akin to exceptions and interrupts
  - sent from the kernel (sometimes at the request of another process) to a process
  - signal type is identified by small integer ID's (1-30)
  - only information in a signal is its ID and the fact that it arrived

| ID | Name    | Default Action   | Corresponding Event                      |
|----|---------|------------------|------------------------------------------|
| 2  | SIGINT  | Terminate        | Interrupt (e.g., dl-c from keyboard)     |
| 9  | SIGKILL | Terminate        | Kill program (cannot override or ignore) |
| 11 | SIGSEGV | Terminate & Dump | Segmentation violation                   |
| 14 | SIGALRM | Terminate        | Timer signal                             |
| 17 | SIGCHLD | Ignore           | Child stopped or terminated              |

## Sending a Signal

- ▶ Kernel **sends** (delivers) a signal to a **destination process** by updating destination process's context
- ▶ Kernel sends a signal for one of the following reasons:
  - Kernel has detected a system event
    - E.g. divide-by-zero (SIGFPE) or child process termination (SIGCHLD)
  - A process calls **kill()** system call to signal to another process
  - Keyboard commands, e.g. ^C, ^Z, etc.

## Receiving a Signal

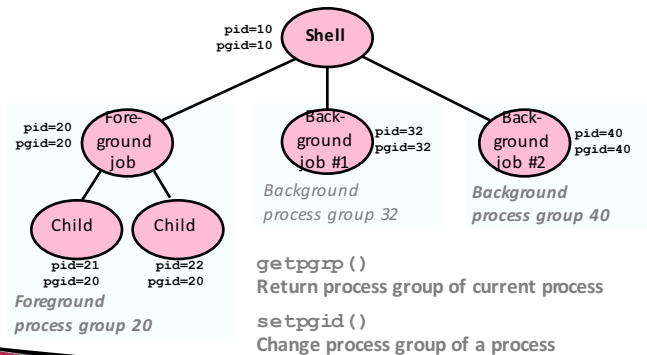
- ▶ A destination process **receives** a signal when kernel forces it to react to the delivery of the signal
- ▶ Three possible ways to react:
  - **Ignore** the signal (do nothing)
  - **Terminate** the process (with optional core dump)
  - **Catch** the signal by executing a user-level **signal handler**
    - Akin to a hardware exception handler being called in response to an asynchronous interrupt

## Signal Concepts

- ▶ Kernel maintains **pending** and **blocked** bit vectors in the context of each process
  - **pending**: represents the set of pending signals
    - Kernel sets bit k in **pending** when a signal of type k is delivered
    - Kernel clears bit k in **pending** when a signal of type k is received
  - **blocked**: represents the set of blocked signals
    - Can be set and cleared by using the **sigprocmask** function

## Process Groups

- Every process belongs to exactly one process group



## Sending Signals with /bin/kill Program

- /bin/kill program sends arbitrary signal to a process or process group

### Examples

- /bin/kill -9 24818  
Send SIGKILL to process 24818

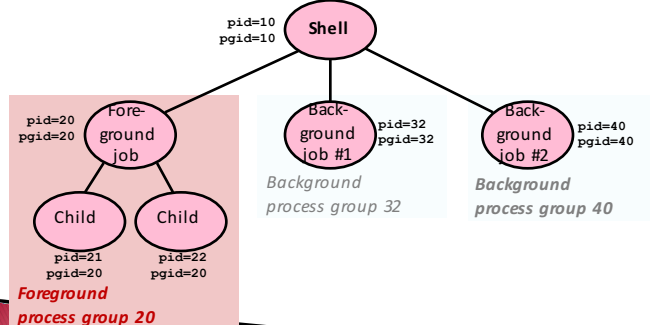
- /bin/kill -9 -24817  
Send SIGKILL to every process in process group 24817

```
linux> ./forks 16
Child1: pid=24818 pgrp=24817
Child2: pid=24819 pgrp=24817
```

```
linux> ps
 PID TTY TIME CMD
 24788 pts/2 00:00:00 tcsh
 24818 pts/2 00:00:02 forks
 24819 pts/2 00:00:02 forks
 24820 pts/2 00:00:00 ps
linux> /bin/kill -9 -24817
linux> ps
 PID TTY TIME CMD
 24788 pts/2 00:00:00 tcsh
 24823 pts/2 00:00:00 ps
linux>
```

## Sending Signals from the Keyboard

- Typing ctrl-c (ctrl-z) sends a SIGINT (SIGTSTP) to every job in the foreground process group.
  - SIGINT – default action is to terminate each process
  - SIGTSTP – default action is to stop (suspend) each process



## Example of ctrl-c and ctrl-z

```
bluefish> ./forks 17
Child: pid=28108 pgrp=28107
Parent: pid=28107 pgrp=28107
```

```
<types ctrl-z>
```

```
Suspended
```

```
bluefish> ps w
```

PID	TTY	STAT	TIME	COMMAND
27699	pts/8	Ss	0:00	-tcsh
28107	pts/8	T	0:01	./forks 17
28108	pts/8	T	0:01	./forks 17
28109	pts/8	R+	0:00	ps w

```
bluefish> fg
```

```
./forks 17
```

```
<types ctrl-c>
```

```
bluefish> ps w
```

PID	TTY	STAT	TIME	COMMAND
27699	pts/8	Ss	0:00	-tcsh
28110	pts/8	R+	0:00	ps w

STAT (process state) Legend:

First letter:

S: sleeping

T: stopped

R: running

Second letter:

s: session leader

+: foreground proc group

See "man ps" for more details

## Sending Signals with kill Function

```
void fork12()
{
 pid_t pid[N];
 int i, child_status;
 for (i = 0; i < N; i++)
 if ((pid[i] = fork()) == 0)
 while(1); /* Child infinite loop */

 /* Parent terminates the child processes */
 for (i = 0; i < N; i++) {
 printf("Killing process %d\n", pid[i]);
 kill(pid[i], SIGINT);
 }

 /* Parent reaps terminated children */
 for (i = 0; i < N; i++) {
 pid_t wpid = wait(&child_status);
 if (WIFEXITED(child_status))
 printf("Child %d terminated with exit status %d\n",
 wpid, WEXITSTATUS(child_status));
 else
 printf("Child %d terminated abnormally\n", wpid);
 }
}
```

## Receiving Signals

- ▶ Suppose kernel is returning from an exception handler and is ready to pass control to process *p*
- ▶ Kernel computes `pnb = pending & ~blocked`
  - The set of pending nonblocked signals for process *p*
- ▶ If (`pnb == 0`)
  - Pass control to next instruction in the logical flow for *p*
- ▶ Else
  - Choose least nonzero bit *k* in `pnb` and force process *p* to **receive** signal *k*
  - The receipt of the signal triggers some **action** by *p*
  - Repeat for all nonzero *k* in `pnb`
  - Pass control to next instruction in logical flow for *p*

## Default Actions

- ▶ Each signal type has a predefined **default action**, which is one of:
  - The process terminates
  - The process terminates and dumps core
  - The process stops until restarted by a SIGCONT signal
  - The process ignores the signal

## Installing Signal Handlers

- ▶ The `signal` function modifies the default action associated with the receipt of signal `signum`:
  - `handler_t *signal(int signum, handler_t *handler)`
- ▶ Different values for `handler`:
  - SIG\_IGN: ignore signals of type `signum`
  - SIG\_DFL: revert to the default action on receipt of signals of type `signum`
  - Otherwise, **handler** is the address of a **signal handler**
    - Called when process receives signal of type `signum`
    - Referred to as **"installing"** the handler
    - Executing handler is called **"catching"** or **"handling"** the signal
    - When the handler executes its return statement, control passes back to instruction in the control flow of the process that was interrupted by receipt of the signal

## Signal Handling Example

```

void int_handler(int sig) {
 safe_printf("Process %d received signal %d\n", getpid(), sig);
 exit(0);
}

void fork13() {
 pid_t pid[N];
 int i, child_status;
 signal(SIGINT, int_handler);
 for (i = 0; i < N; i++) {
 if ((pid[i] = fork()) == 0) {
 while(1); /* child infinite loop */
 }
 for (i = 0; i < N; i++) {
 printf("Killing process %d\n", pid[i]);
 kill(pid[i], SIGINT);
 }
 for (i = 0; i < N; i++) {
 pid_t wpid = wait(&child_status);
 if (WIFEXITED(child_status))
 printf("Child %d terminated with exit status %d\n", wpid, WEXITSTATUS(child_status));
 else
 printf("Child %d terminated with signal %d\n", wpid, WTERMSIG(child_status));
 }
 }
}

```


```

linux> ./forks 13
Killing process 25417
Killing process 25418
Killing process 25419
Killing process 25420
Killing process 25421
Process 25417 received signal 2
Process 25418 received signal 2
Process 25420 received signal 2
Process 25421 received signal 2
Process 25419 received signal 2
Child 25417 terminated with exit status 0
Child 25418 terminated with exit status 0
Child 25420 terminated with exit status 0
Child 25419 terminated with exit status 0
Child 25421 terminated with exit status 0
linux>

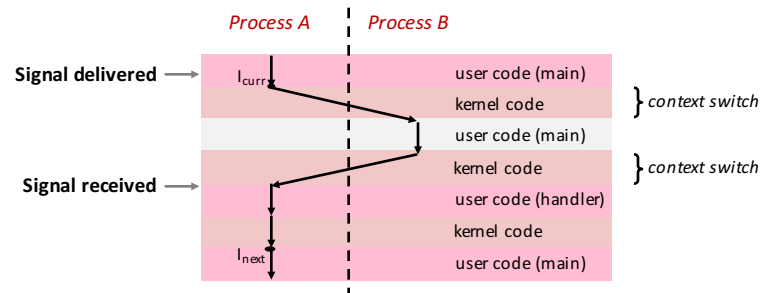
```

## Signals Handlers as Concurrent Flows

- ▶ A signal handler is a separate logical flow (not process) that runs concurrently with the main program
  - “concurrently” in the “not sequential” sense



## Another View of Signal Handlers as Concurrent Flows



## Pending and Blocked Signals

- ▶ A signal is **pending** if sent but not yet received
  - At most one pending signal of any particular type
    - Signals are not queued: If a process has a pending signal of type k, then subsequently sent signals of type k are discarded
- ▶ A process can **block** the receipt of certain signals
  - Blocked signals can be delivered, but will not be received until the signal is unblocked
- ▶ A pending signal is received at most once
- ▶ System calls (and signal handlers) can be interrupted



## Signal Handler Funkiness

```
int ccoun = 0;
void child_handler(int sig)
{
 int child_status;
 pid_t pid = wait(&child_status);
 ccoun--;
 safe_printf(
 "Received signal %d from process %d\n",
 sig, pid);
}

void forks14()
{
 pid_t pid[N];
 int i, child_status;
 ccoun = N;
 signal(SIGCHLD, child_handler);
 for (i = 0; i < N; i++)
 if ((pid[i] = fork()) != 0)
 sleep(1); /* deschi
 exit(0); /* Child
 while (ccoun > 0)
 pause(); /* Suspend until signal occurs */
}
```

▶ Pending signals are not queued  
 ◦ For each signal type, just have single bit indicating whether or not signal is pending  
 ◦ Even if multiple processes have sent this signal

```
linux> ./forks 14
Received SIGCHLD signal 17 for process 21344
Received SIGCHLD signal 17 for process 21345
```

## Living With Nonqueuing Signals

- ▶ Must check for all terminated jobs
  - Typically loop with `wait`

```
void child_handler2(int sig)
{
 int child_status;
 pid_t pid;
 while ((pid = waitpid(-1, &child_status, WNOHANG)) > 0) {
 ccoun--;
 safe_printf("Received signal %d from process %d\n",
 sig, pid);
 }
}

void forks15()
{
 . . .
 signal(SIGCHLD, child_handler2);
 . . .
}
```

```
greatwhite> forks 15
Received signal 17 from process 27476
Received signal 17 from process 27477
Received signal 17 from process 27478
Received signal 17 from process 27479
Received signal 17 from process 27480
greatwhite>
```

## More Signal Handler Funkiness

- ▶ Signal arrival during long system calls (say a `read`)
- ▶ Signal handler interrupts `read` call
  - Linux: upon return from signal handler, the `read` call is restarted automatically
  - Some other flavors of Unix can cause the `read` call to fail with an `EINTR` error number (`errno`)  
in this case, the application program can restart the slow system call
- ▶ Subtle differences like these complicate the writing of portable code that uses signals
  - Consult your textbook for details

## A Program That Reacts to Externally Generated Events (Ctrl-c)

```
#include <stdlib.h>
#include <stdio.h>
#include <signal.h>

void handler(int sig) {
 safe_printf("You think hitting ctrl-c will stop the bomb?\n");
 sleep(2);
 safe_printf("Well...\n");
 sleep(1);
 printf("OK\n");
 exit(0);
}

main() {
 signal(SIGINT, handler); /* inst
 while(1) {
 . . .
 }
}
```

```
linux> ./external
<ctrl-c>
You think hitting ctrl-c will stop
the bomb?
Well...OK
linux>
```

## A Program That Reacts to Internally Generated Events

```
#include <stdio.h>
#include <signal.h>

int beeps = 0;

/* SIGALRM handler */
void handler(int sig) {
 safe_printf("BEEP\n");

 if (++beeps < 5)
 alarm(1);
 else {
 safe_printf("BOOM!\n");
 exit(0);
 }
}
```

```
main() {
 signal(SIGALRM, handler);
 alarm(1); /* send SIGALRM in
 1 second */

 while (1) {
 /* handler returns here */
 }
}
```

```
linux> ./internal
BEEP
BEEP
BEEP
BEEP
BEEP
BOOM!
bass>
```

## Async-Signal-Safety

- Function is *async-signal-safe* if either reentrant (all variables stored on stack frame, CS:APP2e 12.7.2) or non-interruptible by signals.
- Posix guarantees 117 functions to be async-signal-safe
  - `write` is on the list, `printf` is not
- One solution: async-signal-safe wrapper for `printf`:

```
void safe_printf(const char *format, ...) {
 char buf[MAXS];
 va_list args;

 va_start(args, format); /* reentrant */
 vsnprintf(buf, sizeof(buf), format, args); /* reentrant */
 va_end(args); /* reentrant */
 write(1, buf, strlen(buf)); /* async-signal-safe */
}
```

## Today

- ▶ Multitasking, shells
- ▶ Signals
- ▶ Nonlocal jumps

## Nonlocal Jumps: `setjmp/longjmp`

- ▶ Powerful (but dangerous) user-level mechanism for arbitrary control transfer
  - Controlled way to break the procedure call / return discipline
  - Useful for error recovery and signal handling
- ▶ `int setjmp(jmp_buf j)`
  - Must be called before `longjmp`
  - Identifies a return site for a subsequent `longjmp`
  - Called once, returns one or more times
- ▶ Implementation:
  - Remember where you are by storing the current *register context*, *stack pointer*, and *PC value* in `jmp_buf`
  - Return 0

## setjmp/longjmp(cont)

- ▶ `void longjmp(jmp_buf j, int i)`
  - return from the **setjmp** remembered by jump buffer `j` again ...
  - ... this time returning `i` instead of 0
  - Called after **setjmp**
  - Called once, but never returns
- ▶ **longjmp Implementation:**
  - Restore register context (stack pointer, base pointer, PC value) from jump buffer `j`
  - Set `%eax` (the return value) to `i`
  - Jump to the location indicated by the PC stored in jump buf `j`

## Uses for setjmp/longjmp

- ▶ Quickly returning from deep function call nesting
  - E.g. upon error
- ▶ Jump from signal handler to specific code point
- ▶ Checkpoint/restart

## setjmp/longjmp Example

```
#include <setjmp.h>
jmp_buf buf;

main() {
 if (setjmp(buf) != 0) {
 printf("back in main due to an error\n");
 } else {
 printf("first time through\n");
 p1(); /* p1 calls p2, which calls p3 */
 }
 ...
 p3() {
 <error checking code>
 if (error)
 longjmp(buf, 1)
 }
}
```

## Limitations of Nonlocal Jumps


- ▶ Works within stack discipline
  - Can only long jump to environment of function that has been called but not yet completed

```
jmp_buf env;

P1()
{
 if (setjmp(env)) {
 /* Long Jump to here */
 } else {
 P2();
 }
}

P2()
{
 . . . P2(); . . . P3();
}

P3()
{
 longjmp(env, 1);
}
```



## Limitations of Long Jumps (cont.)


- Works within stack discipline
  - Can only long jump to environment of function that has been called but not yet completed

```
jmp_buf env;

P1()
{
 P2(); P3();
}

P2()
{
 if (setjmp(env)) {
 /* Long Jump to here */
 }
}

P3()
{
 longjmp(env, 1);
}
```



## Putting It All Together: A Program That Restarts Itself When `ctrl-c`'d

```
#include <stdio.h>
#include <signal.h>
#include <setjmp.h>

sigjmp_buf buf;

void handler(int sig) {
 siglongjmp(buf, 1);
}

main() {
 signal(SIGINT, handler);

 if (!sigsetjmp(buf, 1))
 printf("starting\n");
 else
 printf("restarting\n");

 while (1) {
 sleep(1);
 printf("processing...\n");
 }
}
```

```
greatwhite> ./restart
starting
processing...
processing...
processing...
restarting ← Ctrl-c
processing...
processing...
restarting ← Ctrl-c
processing...
processing...
```

restart.c