

Linking

CS 341: Intro. to Computer
Architecture & Organization

Andree Jacobson

Today

- Linking
- Case study: Library interpositioning

Example C Program

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

swap.c

```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

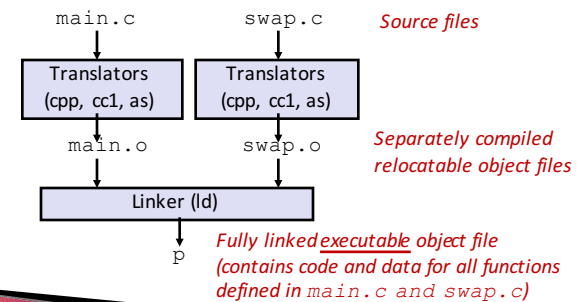
void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

Static Linking

- Programs are translated and linked using a *compiler driver*:

- `unix> gcc -O2 -g -o p main.c swap.c`
- `unix> ./p`



Why Linkers?

► Reason 1: Modularity

- Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library

Why Linkers? (cont)

► Reason 2: Efficiency

- Time: Separate compilation
 - Change one source file, compile, and then relink.
 - No need to recompile other source files.
- Space: Libraries
 - Common functions can be aggregated into a single file...
 - Yet executable files and running memory images contain only code for the functions they actually use.

What Do Linkers Do?

► Step 1. Symbol resolution

- Programs define and reference *symbols* (variables and functions):


```
• void swap() {...}    /* define symbol swap */
• swap();              /* reference symbol a */
• int *xp = &x;        /* define symbol xp, reference x */
```
- Symbol definitions are stored (by compiler) in *symbol table*.
 - Symbol table is an array of structs
 - Each entry includes name, size, and location of symbol.
- Linker maps each symbol reference to **one** symbol definition.

What Do Linkers Do? (cont)

► Step 2. Relocation

- Merges separate code and data sections into single sections
- Relocates symbols from their relative locations in the `.o` files to their final absolute memory locations in the executable.
- Updates all symbol references to reflect their new positions.

Three Kinds of Object Files (Modules)

- ▶ Relocatable object file (.o file)
 - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each .o file is produced from exactly one source (.c) file
- ▶ Executable object file (a.out file)
 - Contains code and data in a form that can be copied directly into memory and then executed.
- ▶ Shared object file (.so file)
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called *Dynamic Link Libraries* (DLLs) by Windows

Executable and Linkable Format (ELF)

- ▶ Standard binary format for object files
- ▶ Originally proposed by AT&T System V Unix
 - Later adopted by BSD Unix variants and Linux
- ▶ One unified format for
 - Relocatable object files (.o),
 - Executable object files (a.out)
 - Shared object files (.so)
- ▶ Generic name: ELF binaries

ELF Object File Format

- ▶ Elf header
 - Word size, byte ordering, file type (.o, exec, so), machine type, etc.
- ▶ Segment header table
 - Page size, virtual addresses memory segments (sections), segment sizes.
- ▶ .text section
 - Code
- ▶ .rodata section
 - Read only data: jump tables, ...
- ▶ .data section
 - Initialized global variables
- ▶ .bss section
 - Uninitialized global variables
 - "Block Started by Symbol"
 - "Better Save Space"
 - Has section header but occupies no space

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

ELF Object File Format (cont.)

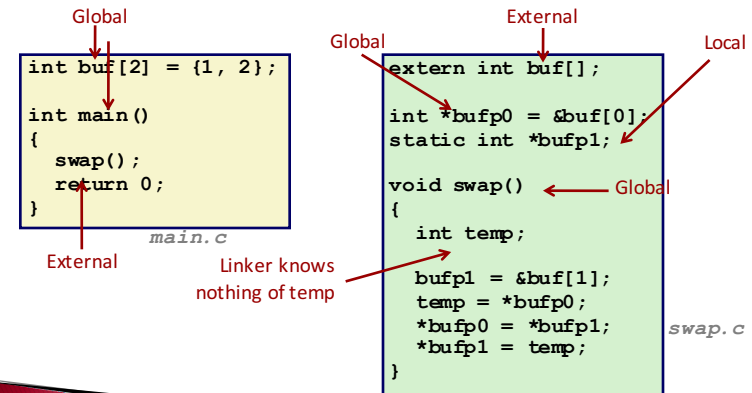
- ▶ .symtab section
 - Symbol table
 - Procedure and static variable names
 - Section names and locations
- ▶ .rel.text section
 - Relocation info for .text section
 - Addresses of instructions that will need to be modified in the executable
 - Instructions for modifying.
- ▶ .rel.data section
 - Relocation info for .data section
 - Addresses of pointer data that will need to be modified in the merged executable
- ▶ .debug section
 - Info for symbolic debugging (gcc -g)
- ▶ Section header table
 - Offsets and sizes of each section

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

Linker Symbols

- ▶ Global symbols
 - Symbols defined by module *m* and visible to other modules.
 - E.g.: non-**static** C functions and non-**static** global variables.
 - External symbols referenced by module *m* but defined by some other module.
- ▶ Local symbols: defined/referenced exclusively by module *m*.
 - E.g.: C functions and variables defined with the **static** attribute.
 - **Local linker symbols are not local program variables**

Resolving Symbols



Symbol Table Entry


```
typedef struct {
    int name;           //string table offset
    int value;          //section offset/VM addr
    int size;           //Object size
    char type:4;        //Data, func, section, fname
    char binding:4;     //Local or global
    char reserved;      //Unused
    char section;       //section header index
}Elf_symbol;
```

Symbol Table Entry

```
typedef struct{
    ELF32_Word      st_name;
    ELF32_Addr      st_value;
    ELF32_Word      st_size;
    unsigned char    st_info;    //type and binding
    unsigned char    st_other;
    ELF32_Half      st_shndx;
} ELF32_Sym;
```

Strong and Weak Symbols

- Program symbols are either strong or weak
 - Strong**: procedures and initialized globals
 - Weak**: uninitialized globals



Linker's Symbol Rules

- Rule 1: Multiple strong symbols are not allowed!
 - Each item can be defined only once
 - Otherwise: Linker error
- Rule 2: Strong symbol and multiple weak symbols? choose the strong symbol!
 - References to the weak symbol resolve to the strong symbol
- Rule 3: Multiple weak symbols? Pick an arbitrary one!
 - Can override this with `gcc -fno-common`

Linker Puzzles

<pre>int x; p1() { }</pre>	<pre>p1() { }</pre>	Link time error: two strong symbols (p1)
<pre>int x; p1() { }</pre>	<pre>int x; p2() { }</pre>	References to <code>x</code> will refer to the same uninitialized int. Is this what you really want?
<pre>int x; int y; p1() { }</pre>	<pre>double x; p2() { }</pre>	Writes to <code>x</code> in <code>p2</code> might overwrite <code>y</code> ! Evil!
<pre>int x=7; int y=5; p1() { }</pre>	<pre>double x; p2() { }</pre>	Writes to <code>x</code> in <code>p2</code> will overwrite <code>y</code> ! Nasty!
<pre>int x=7; p1() { }</pre>	<pre>int x; p2() { }</pre>	References to <code>x</code> will refer to the same initialized variable.

Nightmare scenario: two identical weak structs, compiled by different compilers with different alignment rules.

Role of .h Files

Diagram illustrating the role of header files (.h) in linking:

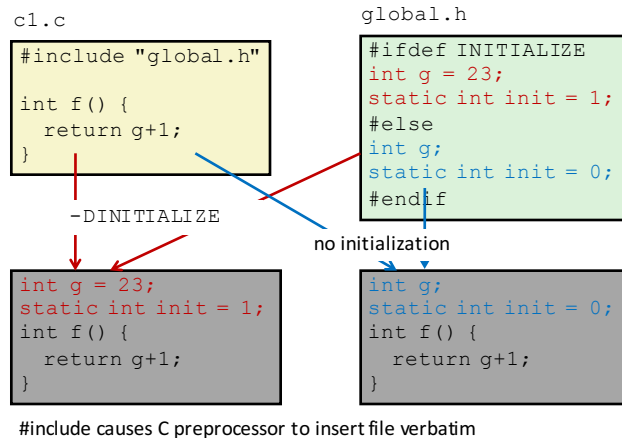
- c1.c**:


```
#include "global.h"
int f() {
    return g+1;
}
```
- c2.c**:

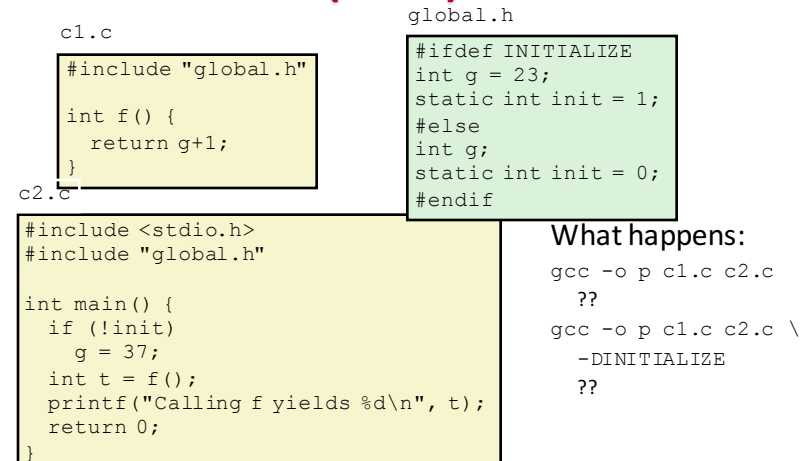

```
#include <stdio.h>
#include "global.h"
int main() {
    if (!init)
        g = 37;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```
- global.h**:


```
#ifndef INITIALIZE
int g = 23;
static int init = 1;
#else
int g;
static int init = 0;
#endif
```

Running Preprocessor



Role of .h Files (cont.)



Global Variables

- ▶ Avoid if you can
- ▶ Otherwise
 - Use **static** if you can
 - Initialize if you define a global variable
 - Use **extern** if you use external global variable

Packaging Commonly Used Functions

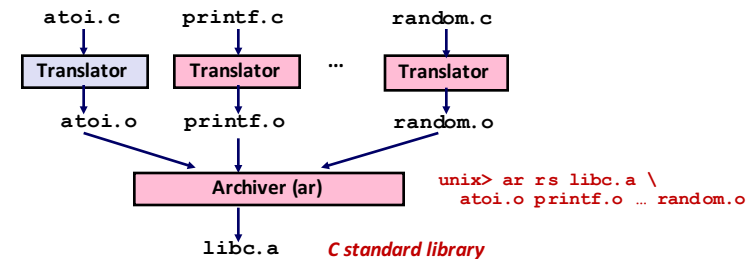
- ▶ How to package commonly used functions ?
 - Math, I/O, memory management, string manipulation, etc.
- ▶ Awkward, given the linker framework so far:
 - **Option 1:** Put all functions into one source file
 - Programmers link big object file into their programs
 - Space and time inefficient
 - **Option 2:** Put each function in a separate source file
 - Programmers explicitly link appropriate object files
 - More efficient, but burdensome on the programmer

Solution: Static Libraries

► Static libraries (.a archive files)

- Concatenate related relocatable object files into a single file with an index (called an *archive*).
- Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
- If an archive member file resolves reference, link it into the executable.

Creating Static Libraries



- Archiver allows incremental updates
- Recompile function that changes and replace .o file in archive.

Commonly Used Libraries

libc.a (the C standard library)

- 8 MB archive of 1392 object files.
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

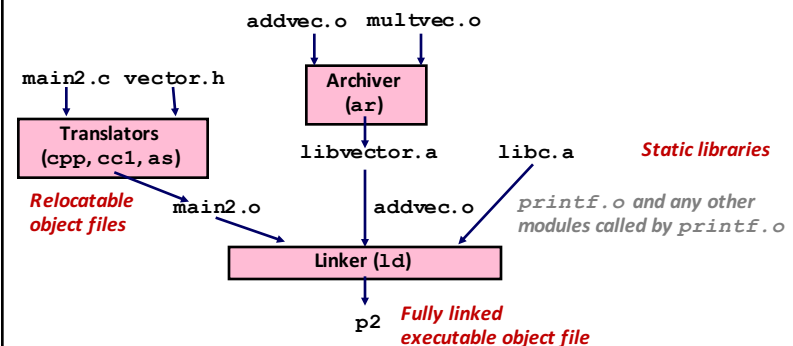
libm.a (the C math library)

- 1 MB archive of 401 object files.
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t /usr/lib/libc.a | sort
...
fork.o
...
fprintf.o
fpu_control.o
putc.o
freopen.o
fscanf.o
fseek.o
fstab.o
...
```

```
% ar -t /usr/lib/libm.a | sort
...
e_acos.o
e_acosf.o
e_acosh.o
e_acoshf.o
e_acoshl.o
e_acosl.o
e_asin.o
e_asinf.o
e_asinl.o
...
```

Linking with Static Libraries



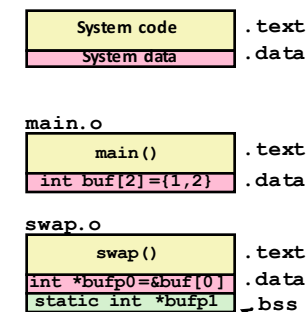
Using Static Libraries

- ▶ Linker's algorithm for resolving external references:
 1. Scan .o files and .a files in the command line order.
 2. During scan, keep a list of the current unresolved references.
 3. As each new .o or .a file, *obj*, is encountered, try to resolve each unresolved reference in the list against the symbols defined in *obj*.
 4. If any entries in the unresolved list at end of scan, then error.
- ▶ Problem:
 - Command line order matters!
 - Moral: put libraries at the end of the command line.

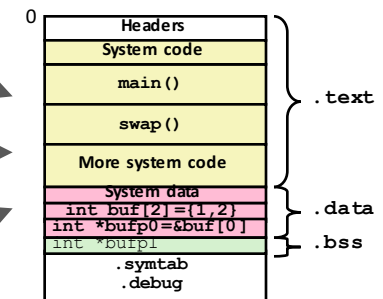
```
unix> gcc -L. libtest.o -lmine
unix> gcc -L. -lmine libtest.o
libtest.o: In function `main':
libtest.o(.text+0x4): undefined reference to `libfun'
```

Relocating Code and Data

Relocatable Object Files



Executable Object File



Even though private to swap, requires allocation in .bss

Relocation Info (main)

main.c

```
int buf[2] = {1,2};

int main()
{
    swap();
    return 0;
}
```

main.o

```
00000000 <main>:
0: 8d 4c 24 04      lea    0x4(%esp),%ecx
4: 83 e4 f0         and    $0xfffffffff0,%esp
7: ff 71 fc         pushl  0xfffffffffc(%ecx)
a: 55              push   %ebp
b: 89 e5           mov    %esp,%ebp
d: 51             push   %ecx
e: 83 ec 04        sub    $0x4,%esp
11: e8 fc ff ff     call   12<main+0x12>
12: R_386_PC32 swap
16: 83 c4 04        add    $0x4,%esp
19: 31 c0           xor    %eax,%eax
1b: 59             pop    %ecx
1c: 5d             pop    %ebp
1d: 8d 61 fc        lea    0xfffffffffc(%ecx),%esp
20: c3             ret
```

Disassembly of section .data:

```
00000000 <buf>:
0: 01 00 00 00 02 00 00 00
```

Source: objdump -r -d

Relocation Info (swap, .text)

swap.c

```
extern int buf[];

int
*bufp0 = &buf[0];

static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

swap.o

Disassembly of section .text:

```
00000000 <swap>:
0: 8b 15 00 00 00 00      mov    0x0,%edx
2: R_386_32 buf
6: a1 04 00 00 00 00      mov    0x4,%eax
7: R_386_32 buf
b: 55              push   %ebp
c: 89 e5           mov    %esp,%ebp
e: c7 05 00 00 00 00 04    movl   $0x4,0x0
15: 00 00 00
10: R_386_32 .bss
14: R_386_32 buf
18: 8b 08           mov    (%eax),%ecx
1a: 89 10           mov    %edx,(%eax)
1c: 5d             pop    %ebp
1d: 89 0d 04 00 00 00      mov    %ecx,0x4
1f: R_386_32 buf
23: c3             ret
```


Relocation Info (swap, .data)

swap.c

```
extern int buf[];

int *bufp0 =
&buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

Disassembly of section .data:

```
00000000 <bufp0>:
0: 00 00 00 00

0: R_386_32 buf
```

Executable Before/After Relocation (.text)

00000000 <main>:

```
e: 83 ec 04 sub $0x4,%esp
11: e8 fc ff ff call 12 <main+0x12>
12: R_386_PC32 swap
16: 83 c4 04 add $0x4,%esp
```

0x8048396 + 0x1a
= 0x80483b0

08048380 <main>:

```
8048380: 8d 4c 24 04 lea 0x4(%esp),%ecx
8048384: 83 e4 f0 and $0xfffffffff0,%esp
8048387: ff 71 fc pushl 0xfffffffff(%ecx)
804838a: 55 push %ebp
804838b: 89 e5 mov %esp,%ebp
804838d: 51 push %ecx
804838e: 83 ec 04 sub $0x4,%esp
8048391: e8 1a 00 00 00 call 80483b0 <swap>
8048396: 83 c4 04 add $0x4,%esp
8048399: 31 c0 xor %eax,%eax
804839b: 59 pop %ecx
804839c: 5d pop %ebp
804839d: 8d 61 fc lea 0xfffffffff(%ecx),%esp
80483a0: c3 ret
```

```
0: 8b 15 00 00 00 00 mov 0x0,%edx
2: R_386_32 buf
6: a1 04 00 00 00 00 mov 0x4,%eax
7: R_386_32 buf
...
e: c7 05 00 00 00 00 04 movl $0x4,0x0
15: 00 00 00
10: R_386_32 .bss
14: R_386_32 buf
1d: 89 0d 04 00 00 00 mov %ecx,0x4
1f: R_386_32 buf
23: c3 ret
```

```
080483b0 <swap>:
80483b0: 8b 15 20 96 04 08 mov 0x8049620,%edx
80483b6: a1 24 96 04 08 mov 0x8049624,%eax
80483bb: 55 push %ebp
80483bc: 89 e5 mov %esp,%ebp
80483be: c7 05 30 96 04 08 24 movl $0x8049624,0x8049630
80483c5: 96 04 08
80483c8: 8b 08 mov (%eax),%ecx
80483ca: 89 10 mov %edx,(%eax)
80483cc: 5d pop %ebp
80483cd: 89 0d 24 96 04 08 mov %ecx,0x8049624
80483d3: c3 ret
```

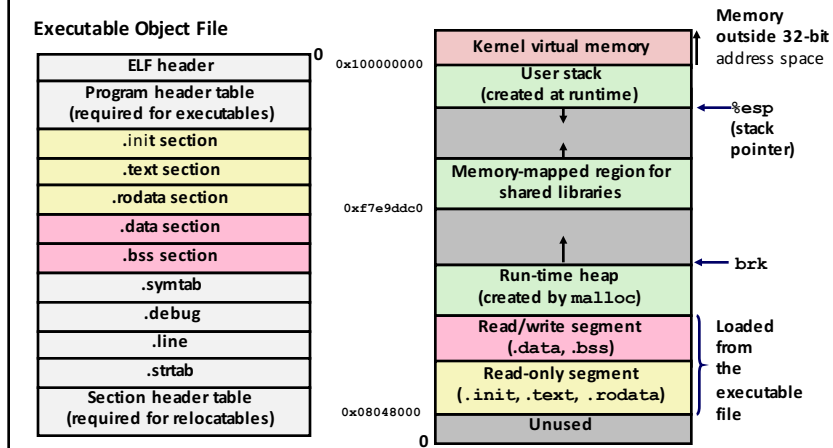
Executable After Relocation (.data)

Disassembly of section .data:

```
08049620 <buf>:
8049620: 01 00 00 00 02 00 00 00

08049628 <bufp0>:
8049628: 20 96 04 08
```

Loading Executable Object Files



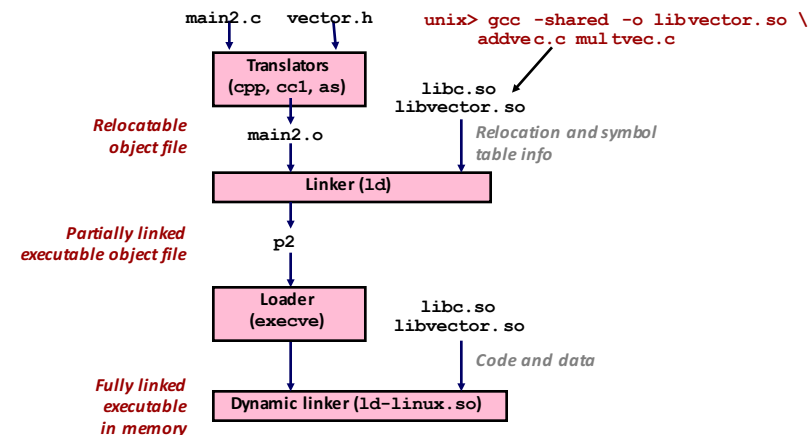
Shared Libraries

- ▶ Static libraries have the following disadvantages:
 - Duplication in the stored executables (every function need std libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to explicitly relink
- ▶ Modern solution: Shared Libraries
 - Object files that contain code and data that are loaded and linked into an application *dynamically*, at either *load-time* or *run-time*
 - Also called: dynamic link libraries, DLLs, .so files

Shared Libraries (cont.)

- ▶ Dynamic linking can occur when executable is first loaded and run (load-time linking).
 - Common case for Linux, handled automatically by the dynamic linker (**ld-linux.so**).
 - Standard C library (**libc.so**) usually dynamically linked.
- ▶ Dynamic linking can also occur after program has begun (run-time linking).
 - In Linux, this is done by calls to the **dlopen()** interface.
 - Distributing software.
 - High-performance web servers.
 - Runtime library interpositioning.
- ▶ Shared library routines can be shared by multiple processes.

Dynamic Linking at Load-time



Dynamic Linking at Run-time

```
#include <stdio.h>
#include <dlfcn.h>

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main()
{
    void *handle;
    void (*addvec)(int *, int *, int *, int);
    char *error;

    /* dynamically load the shared lib that contains addvec() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
}
```

Dynamic Linking at Run-time

```
...

/* get a pointer to the addvec() function we just loaded */
addvec = dlsym(handle, "addvec");
if ((error = dlerror()) != NULL) {
    fprintf(stderr, "%s\n", error);
    exit(1);
}

/* Now we can call addvec() just like any other function */
addvec(x, y, z, 2);
printf("z = [%d %d]\n", z[0], z[1]);

/* unload the shared library */
if (dlclose(handle) < 0) {
    fprintf(stderr, "%s\n", dlerror());
    exit(1);
}

return 0;
}
```

Today

- Linking
- Case study: Library interpositioning

Case Study: Library Interpositioning

- Library interpositioning : powerful linking technique for intercepting arbitrary function calls
- Interpositioning can occur at:
 - Compile time: When the source code is compiled
 - Link time: When the relocatable object files are statically linked to form an executable object file
 - Load/run time: When an executable object file is loaded into memory, dynamically linked, and then executed.

Some Interpositioning Applications

- ▶ Security
 - Confinement (sandboxing)
 - Interpose calls to libc functions.
 - Behind the scenes encryption
 - Automatically encrypt otherwise unencrypted network connections.
- ▶ Monitoring and Profiling
 - Count number of calls to functions
 - Characterize call sites and arguments to functions
 - Malloc tracing
 - Detecting memory leaks
 - **Generating address traces**

Example program

```
#include <stdio.h>
#include <stdlib.h>
#include <malloc.h>

int main()
{
    free(malloc(10));
    printf("hello, world\n");
    exit(0);
}
```

hello.c

- ▶ Goal: trace the addresses and sizes of the allocated and freed blocks, without modifying the source code
- ▶ Three solutions: interpose on the lib malloc and free functions at compile time, link time, and load/run time.

Compile-time Interpositioning

```
#ifndef COMPILETIME
/* Compile-time interposition of malloc and free using C
 * preprocessor. A local malloc.h file defines malloc (free)
 * as wrappers mymalloc (myfree) respectively.
 */

#include <stdio.h>
#include <malloc.h>

/*
 * mymalloc - malloc wrapper function
 */
void *mymalloc(size_t size, char *file, int line)
{
    void *ptr = malloc(size);
    printf("%s:%d: malloc(%d)=%p\n", file, line, (int)size, ptr);
    return ptr;
}
```

mymalloc.c

Compile-time Interpositioning

```
#define malloc(size) mymalloc(size, __FILE__, __LINE__)
#define free(ptr) myfree(ptr, __FILE__, __LINE__)

void *mymalloc(size_t size, char *file, int line);
void myfree(void *ptr, char *file, int line);
```

malloc.h

```
linux> make helloc
gcc -O2 -Wall -DCOMPILETIME -c mymalloc.c
gcc -O2 -Wall -I. -o helloc hello.c mymalloc.o
linux> make runc
./helloc
hello.c:7: malloc(10)=0x501010
hello.c:7: free(0x501010)
hello, world
```

Link-time Interpositioning

```
#ifndef LINKTIME
/* Link-time interposition of malloc and free using the static
linker's (ld) "--wrap symbol" flag. */

#include <stdio.h>

void *__real_malloc(size_t size);
void __real_free(void *ptr);

/*
 * __wrap_malloc - malloc wrapper function
 */
void *__wrap_malloc(size_t size)
{
    void *ptr = __real_malloc(size);
    printf("malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}

```

mymalloc.c

Link-time Interpositioning

```
linux> make hellol
gcc -O2 -Wall -DLINKTIME -c mymalloc.c
gcc -O2 -Wall -Wl,--wrap,malloc -Wl,--wrap,free \
-o hellol hello.c mymalloc.o
linux> make runl
./hellol
malloc(10) = 0x501010
free(0x501010)
hello, world

```

- ▶ The “-Wl” flag passes argument to linker
- ▶ “--wrap, malloc” tells linker to resolve:
 - Refs to malloc should be resolved as __wrap_malloc
 - Refs to __real_malloc should be resolved as malloc

Load/Run-time Interpositioning

```
#ifndef RUNTIME
/* Run-time interposition of malloc and free based on
 * dynamic linker's (ld-linux.so) LD_PRELOAD mechanism */
#define GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <dlfcn.h>

void *malloc(size_t size)
{
    static void *(*mallocp)(size_t size);
    char *error;
    void *ptr;

    /* get address of libc malloc */
    if (!mallocp) {
        mallocp = dlsym(RTLD_NEXT, "malloc");
        if ((error = dlerror()) != NULL) {
            fputs(error, stderr);
            exit(1);
        }
    }
    ptr = mallocp(size);
    printf("malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}

```

mymalloc.c

Load/Run-time Interpositioning

```
linux> make hellor
gcc -O2 -Wall -DRUNTIME -shared -fPIC -o mymalloc.so mymalloc.c
gcc -O2 -Wall -o hellor hello.c
linux> make runr
(LD_PRELOAD="/usr/lib64/libdl.so ./mymalloc.so" ./hellor)
malloc(10) = 0x501010
free(0x501010)
hello, world

```

- ▶ LD_PRELOAD environment variable tells the dynamic linker to resolve unresolved refs (e.g., to malloc) by looking in libdl.so and mymalloc.so first.
 - libdl.so necessary to resolve references to the dlopen functions.

Interpositioning Recap

- ▶ Compile Time
 - Apparent calls to malloc/free get macro-expanded into calls to mymalloc/myfree
- ▶ Link Time
 - Use linker trick to have special name resolutions
 - malloc → __wrap_malloc
 - __real_malloc → malloc
- ▶ Compile Time
 - Implement custom version of malloc/free that use dynamic linking to load library malloc/free under different names