

Virtual Memory: Concepts


CS 341: Intro. to Computer
Architecture & Organization

Andree Jacobson

Today

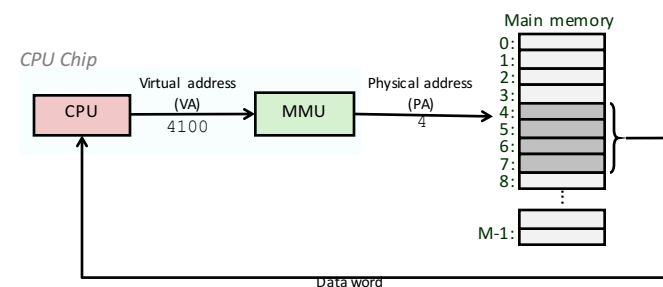
- Address spaces
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation

A System Using Physical Addressing



- Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and

A System Using Virtual Addressing



- Used in all modern servers, desktops, and laptops
- One of the great ideas in computer science

Address Spaces

- ▶ **Linear address space:** Ordered set of contiguous non-negative integer addresses:
 $\{0, 1, 2, 3 \dots\}$
- ▶ **Virtual address space:** Set of $N = 2^n$ virtual addresses
 $\{0, 1, 2, 3, \dots, N-1\}$
- ▶ **Physical address space:** Set of $M = 2^m$ physical addresses
 $\{0, 1, 2, 3, \dots, M-1\}$
- ▶ Clean distinction between data (bytes) and their attributes (addresses)
- ▶ Each object can now have multiple addresses
- ▶ Every byte in main memory:
 one physical address, one (or more) virtual addresses

Why Virtual Memory (VM)?


- ▶ Uses main memory efficiently
 - Use DRAM as a cache for the parts of a virtual address space
- ▶ Simplifies memory management
 - Each process gets the same uniform linear address space
- ▶ Isolates address spaces
 - One process can't interfere with another's memory
 - User program cannot access privileged kernel information

Today

- ▶ Address spaces
- ▶ VM as a tool for caching
- ▶ VM as a tool for memory management
- ▶ VM as a tool for memory protection
- ▶ Address translation

VM as a Tool for Caching

- ▶ **Virtual memory** is an array of N contiguous bytes on disk
- ▶ The contents of the array on disk are cached in **physical memory (DRAM cache)**
 - These cache blocks are called *pages* (size is $P = 2^p$ bytes)

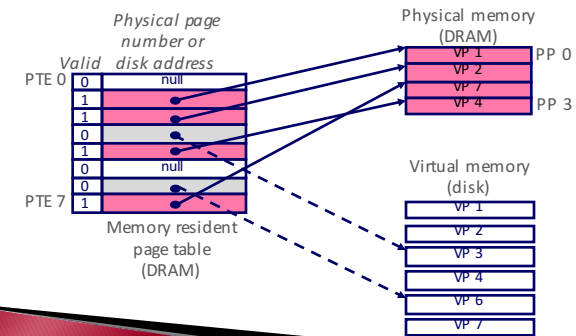


DRAM Cache Organization

- ▶ DRAM cache organization driven by the enormous miss penalty
 - DRAM is about **10x** slower than SRAM
 - Disk is about **10,000x** slower than DRAM
- ▶ Consequences
 - Large page (block) size: typically 4-8 KB, sometimes 4 MB
 - Fully associative
 - Any VP can be placed in any PP
 - Requires a “large” mapping function – different from CPU caches
 - Highly sophisticated, expensive replacement algorithms
 - Too complicated and open-ended to be implemented in hardware
 - Write-back rather than write-through


Page Tables

- ▶ A **page table** is an array of page table entries (PTEs) that maps virtual pages to physical pages.
 - Per-process kernel data structure in DRAM




Page Hit

- ▶ **Page hit**: reference to VM word that is in physical memory (DRAM cache hit)




Page Fault

- ▶ **Page fault**: reference to VM word that is not in physical memory (DRAM cache miss)




Handling Page Fault

- Page miss causes page fault (an exception)




Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)




Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)



Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending instruction is restarted: page hit!



Locality to the Rescue Again!


- ▶ Virtual memory works because of locality
- ▶ At any point in time, programs tend to access a set of active virtual pages called the *working set*
 - Programs with better temporal locality will have smaller working sets
- ▶ If (working set size < main memory size)
 - Good performance for one process after compulsory misses
- ▶ If (SUM(working set sizes) > main memory size)
 - *Thrashing*: Performance meltdown where pages are swapped (copied) in and out continuously

Today

- ▶ Address spaces
- ▶ VM as a tool for caching
- ▶ VM as a tool for memory management
- ▶ VM as a tool for memory protection
- ▶ Address translation


VM as a Tool for Memory Management

- ▶ Key idea: each process has its own virtual address space
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory
 - Well chosen mappings simplify memory allocation and management



VM as a Tool for Memory Management

- ▶ Memory allocation
 - Each virtual page can be mapped to any physical page
 - A virtual page can be stored in different physical pages at different times
- ▶ Sharing code and data among processes
 - Map virtual pages to the same physical page (here: PP 6)



Today

- Address spaces
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation

VM as a Tool for Memory Protection

- Extend PTEs with permission bits
- Page fault handler checks these before remapping
 - If violated, send process SIGSEGV (segmentation fault)

