

Memory Allocation

CS 341: Intro. to Computer
Architecture & Organization

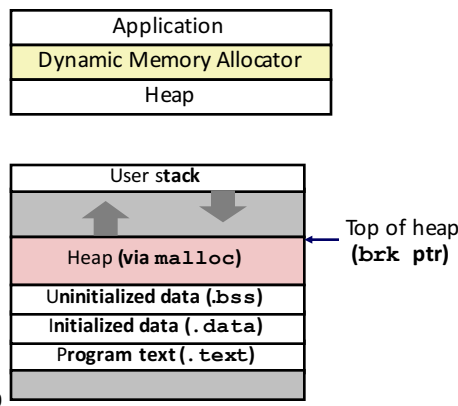
Andree Jacobson

Today

- Basic concepts
- Implicit free lists

Dynamic Memory Allocation

- Programmers use *dynamic memory allocators* (such as `malloc`) to acquire VM at run time.
 - For data structures whose size is only known at runtime.
- Dynamic memory allocators manage an area of process virtual memory known as the *heap*.



Dynamic Memory Allocation

- Allocator maintains heap as collection of variable sized *blocks*, which are either *allocated* or *free*
- Types of allocators
 - **Explicit allocator:** application allocates and frees space
 - E.g., `malloc` and `free` in C
 - **Implicit allocator:** application allocates, but does not free space
 - E.g. garbage collection in Java, ML, and Lisp
- Will discuss simple explicit memory allocation today

The malloc Package

```
#include <stdlib.h>
```

```
void *malloc(size_t size)
```

- Successful:
 - Returns a pointer to a memory block of at least **size** bytes (typically) aligned to 8-byte boundary
 - If **size == 0**, returns NULL
- Unsuccessful: returns NULL (0) and sets **errno**

```
void free(void *p)
```

- Returns the block pointed at by **p** to pool of available memory
- **p** must come from a previous call to **malloc** or **realloc**

Other functions

- **calloc**: Version of **malloc** that initializes allocated block to zero.
- **realloc**: Changes the size of a previously allocated block.
- **sbrk**: Used internally by allocators to grow or shrink the heap

malloc Example

```
void foo(int n, int m) {
    int i, *p;


    /* Allocate a block of n ints */
    p = (int *) malloc(n * sizeof(int));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }

    /* Initialize allocated block */
    for (i=0; i<n; i++)
        p[i] = i;

    /* Return p to the heap */
    free(p);
}
```

Assumptions Made in This Lecture

- Memory is word addressed (each word can hold a pointer)



Allocation Example



Constraints

- ▶ Applications
 - Can issue arbitrary sequence of **malloc** and **free** requests
 - **free** request must be to a **malloc'd** block
- ▶ Allocators
 - Can't control number or size of allocated blocks
 - Must respond immediately to **malloc** requests
 - *i.e.*, can't reorder or buffer requests
 - Must allocate blocks from free memory
 - *i.e.*, can only place allocated blocks in free memory
 - Must align blocks so they satisfy all alignment requirements
 - 8 byte alignment for GNU **malloc** (**libc malloc**) on Linux boxes
 - Can manipulate and modify only free memory
 - Can't move the allocated blocks once they are **malloc'd**
 - *i.e.*, compaction is not allowed

Performance Goal: Throughput

- ▶ Given some sequence of **malloc** and **free** requests:
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- ▶ Goals: maximize throughput and peak memory utilization
 - These goals are often conflicting
- ▶ Throughput:
 - Number of completed requests per unit time
 - Example:
 - 5,000 **malloc** calls and 5,000 **free** calls in 10 seconds
 - Throughput is 1,000 operations/second

Performance Goal: Peak Memory Utilization


- ▶ Given some sequence of **malloc** and **free** requests:
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- ▶ **Def:** Aggregate payload P_k
 - **malloc** (**p**) results in a block with a **payload** of **p** bytes
 - After request R_k has completed, the **aggregate payload** P_k is the sum of currently allocated payloads
- ▶ **Def:** Current heap size H_k
 - Assume H_k is monotonically nondecreasing
 - *i.e.*, heap only grows when allocator uses **sbrk**
- ▶ **Def:** Peak memory utilization after k requests
 - $U_k = (\max_{i \leq k} P_i) / H_k$

Fragmentation

- ▶ Poor memory utilization caused by **fragmentation**
 - **internal** fragmentation
 - **external** fragmentation

Internal Fragmentation

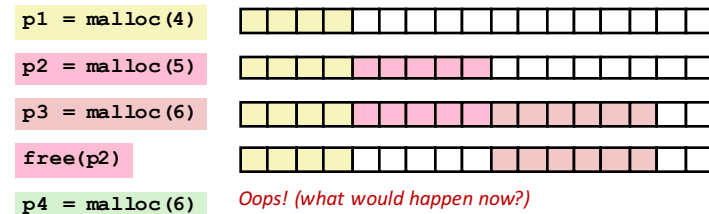
- For a given block, *internal fragmentation* occurs if payload is smaller than block size



- Caused by
 - Overhead of maintaining heap data structures
 - Padding for alignment purposes
 - Explicit policy decisions (e.g., to return a big block to satisfy a small request)
- Depends only on the pattern of *previous* requests
 - Thus, easy to measure

External Fragmentation

- Occurs when there is enough aggregate heap memory, but no single free block is large enough



- Depends on the pattern of future requests
 - Thus, difficult to measure

Implementation Issues

- How do we know how much memory to free given just a pointer?
- How do we keep track of the free blocks?
- What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?
- How do we pick a block to use for allocation -- many might fit?
- How do we reinsert freed block?

Knowing How Much to Free

- Standard method
 - Keep the length of a block in the word preceding the block.
 - This word is often called the *header field* or *header*
 - Requires an extra word for every allocated block



Keeping Track of Free Blocks

- ▶ Method 1: *Implicit list* using length—links all blocks



- ▶ Method 2: *Explicit list* among the free blocks using pointers




- ▶ Method 3: *Segregated free list*
 - Different free lists for different size classes
- ▶ Method 4: *Blocks sorted by size*
 - Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Today


- ▶ Basic concepts
- ▶ Implicit free lists

Method 1: Implicit List

- ▶ For each block we need both size and allocation status
 - Could store this information in two words: wasteful!
- ▶ Standard trick
 - If blocks are aligned, some low-order address bits are always 0
 - Instead of storing an always-0 bit, use it as an allocated/free flag
 - When reading size word, must mask out this bit



Detailed Implicit Free List Example



Implicit List: Finding a Free Block

► First fit:

- Search list from beginning, choose **first** free block that fits:

```
p = start;
while ((p < end) &&      // not passed end
      ((*p & 1) ||      // already allocated
      (*p <= len)))     // too small
    p = p + (*p & -2);   // goto next block (word addressed)
```

- Can take linear time in total number of blocks (allocated and free)
- In practice it can cause “splinters” at beginning of list

► Next fit:

- Like first fit, but search list starting where previous search finished
- Should often be faster than first fit: avoids re-scanning unhelpful blocks
- Some research suggests that fragmentation is worse


► Best fit:

- Search the list, choose the **best** free block: fits, with fewest bytes left over
- Keeps fragments small—usually helps fragmentation
- Will typically run slower than first fit

Implicit List: Allocating in Free Block

► Allocating in a free block: *splitting*

- Since allocated space might be smaller than free space, we might want to split the block



addblock(p, 4)



```
void addblock(ptr p, int len) {
    int newsize = ((len + 1) >> 1) << 1; // round up to even
    int oldsize = *p & -2;                // mask out low bit
    *p = newsize | 1;                     // set new length
    if (newsize < oldsize)
        *(p+newsize) = oldsize - newsize; // set length in remaining
                                           // part of block
}
```


Implicit List: Freeing a Block

► Simplest implementation:

- Need only clear the “allocated” flag

```
void free_block(ptr p) { *p = *p & -2 }
```

- But can lead to “false fragmentation”




There is enough free space, but the allocator won't be able to find it

Implicit List: Coalescing

► Join (*coalesce*) with next/previous blocks, if free

- Coalescing with next block




```
void free_block(ptr p) {
    *p = *p & -2; // clear allocated flag
    next = p + *p; // find next block
    if ((*next & 1) == 0)
        *p = *p + *next; // add to this block if
                          // not allocated
}
```

- But how do we coalesce with *previous* block?


Implicit List: Bidirectional Coalescing

► Boundary tags [Knuth73]


- Replicate size/allocated word at “bottom” (end) of free blocks
- Allows us to traverse the “list” backwards, but requires extra space
- Important and general technique!




*Format of
allocated and
free blocks*




Constant Time Coalescing




Constant Time Coalescing (Case 1)




Constant Time Coalescing (Case 2)



Constant Time Coalescing (Case 3)



Constant Time Coalescing (Case 4)



Disadvantages of Boundary Tags

- ▶ Internal fragmentation
- ▶ Can it be optimized?
 - Which blocks need the footer tag?
 - What does that mean?

Summary of Key Allocator Policies

- ▶ Placement policy:
 - First-fit, next-fit, best-fit, etc.
 - Trades off lower throughput for less fragmentation
 - **Interesting observation:** segregated free lists approximate a best fit placement policy without having to search entire free list
- ▶ Splitting policy:
 - When do we go ahead and split free blocks?
 - How much internal fragmentation are we willing to tolerate?
- ▶ Coalescing policy:
 - **Immediate coalescing:** coalesce each time **free** is called
 - **Deferred coalescing:** try to improve performance of **free** by deferring coalescing until needed. Examples:
 - Coalesce as you scan the free list for **malloc**
 - Coalesce when external fragmentation reaches some threshold

Implicit Lists: Summary

- ▶ Implementation: very simple
- ▶ Allocate cost:
 - linear time worst case
- ▶ Free cost:
 - constant time worst case
 - even with coalescing
- ▶ Memory usage:
 - will depend on placement policy
 - First-fit, next-fit or best-fit
- ▶ Not used in practice for `malloc/free` because of linear-time allocation
 - used in many special purpose applications
- ▶ However, the concepts of splitting and boundary tag coalescing are general to *all* allocators