

Concurrent Programming

CS 341: Intro. to Computer
Architecture & Organization

Andree Jacobson

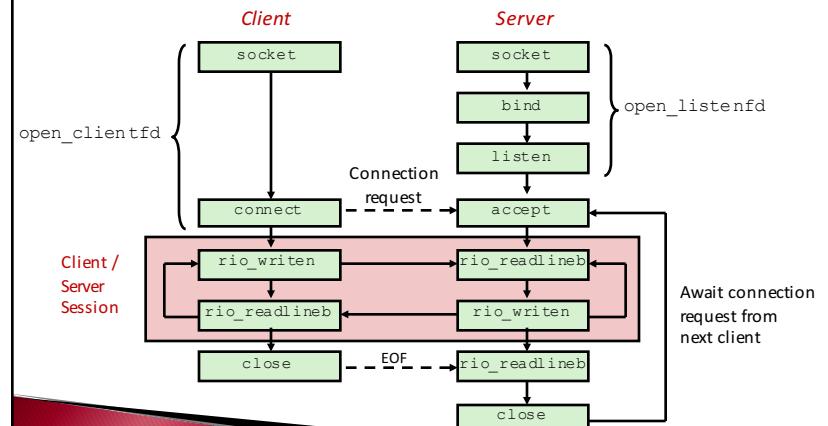
Concurrent Programming is Hard!

- ▶ The human mind tends to be (consciously) sequential
- ▶ The notion of time is often misleading
- ▶ We can't imagine complete sequence of computer events

Concurrent Programming is Hard!


- ▶ Classical problem classes of concurrent programs:
 - **Races**: outcome depends on arbitrary scheduling decisions elsewhere in the system
 - Example: musical chairs ☺
 - **Deadlock**: improper resource allocation prevents forward progress
 - Example: traffic gridlock
 - **Livelock / Starvation / Fairness**: external events and/or system scheduling decisions can prevent sub-task progress
 - Example: people always jump in front of you in line

Iterative Echo Server




Iterative Servers

- Iterative servers process one request at a time




Where Does Second Client Block?

- Second client attempts to connect to iterative server



- Call to connect returns
 - Even though connection not yet accepted
 - Server side TCP manager queues request
 - Feature known as "TCP listen backlog"
- Call to rio_writen returns
 - Server side TCP manager buffers input data
- Call to rio_readlineb blocks
 - Server hasn't written anything for it to read yet.

Fundamental Flaw of Iterative Servers



- Solution: use *concurrent servers* instead
 - Concurrent servers use concurrent flows to serve concurrent clients


Creating Concurrent Flows

Allow server to handle multiple clients simultaneously

- 1. Processes
 - Kernel automatically interleaves multiple logical flows
 - Each flow has its own private address space
- 2. Threads
 - Kernel automatically interleaves multiple logical flows
 - Each flow shares the same address space
- 3. I/O multiplexing with `select()`
 - Programmer manually interleaves multiple logical flows
 - All flows share the same address space
 - Relies on lower-level system abstractions

Concurrent Servers: Multiple Processes

- Spawn separate process for each client



Review: Iterative Echo Server

```

int main(int argc, char **argv)
{
    int listenfd, connfd;
    int port = atoi(argv[1]);
    struct sockaddr_in clientaddr;
    int clientlen = sizeof(clientaddr);

    listenfd = Open_listenfd(port);
    while (1) {
        connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
        echo(connfd);
        Close(connfd);
    }
    exit(0);
}
  
```

- Accept a connection request
- Handle echo requests until client terminates

Process-Based Concurrent Server

```

int main(int argc, char **argv)
{
    int listenfd, connfd;
    int port = atoi(argv[1]);
    struct sockaddr_in clientaddr;
    int clientlen = sizeof(clientaddr);

    Signal(SIGCHLD, sigchld_handler);
    listenfd = Open_listenfd(port);
    while (1) {
        connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
        if (Fork() == 0) {
            Close(listenfd); /* Child closes its listening socket */
            echo(connfd); /* Child services client */
            Close(connfd); /* Child closes connection with client */
            exit(0); /* Child exits */
        }
        Close(connfd); /* Parent closes connected socket (important!) */
    }
}
  
```

Fork separate process for each client
Does not allow any communication between different client handlers

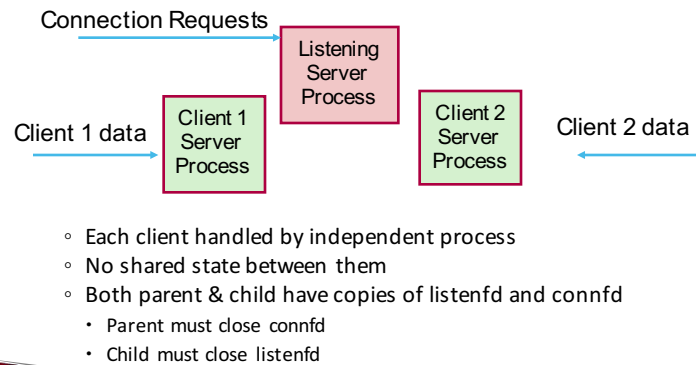
Process-Based Concurrent Server (cont)

```

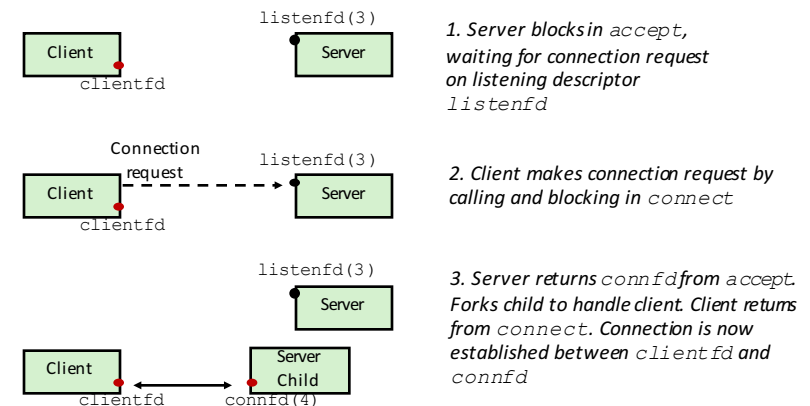
void sigchld_handler(int sig)
{
    while (waitpid(-1, 0, WNOHANG) > 0)
        ;
    return;
}
  
```

- Reap all zombie children

Process Execution Model



Concurrent Server: `accept` Illustrated



Implementation Must-dos With Process-Based Designs

- ▶ Listening server process must reap zombie children
 - to avoid fatal memory leaks
- ▶ Listening server process must close its copy of `connfd`
 - Kernel keeps reference for each socket/open file
 - After fork, `refcnt(connfd) = 2`
 - Connection will not be closed until `refcnt(connfd) == 0`

Pros and Cons of Process-Based Designs

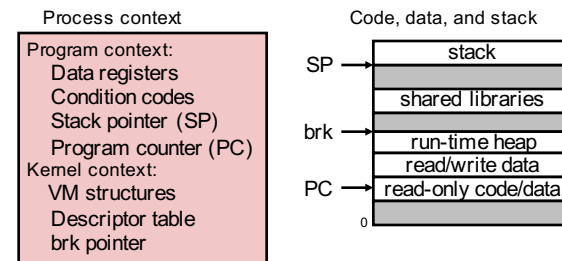
- ▶ + Handle multiple connections concurrently
- ▶ + Clean sharing model
 - descriptors (no)
 - file tables (yes)
 - global variables (no)
- ▶ + Simple and straightforward
- ▶ – Additional overhead for process control
- ▶ – Nontrivial to share data between processes
 - Requires IPC (interprocess communication) mechanisms
 - FIFO's (named pipes), System V shared memory and semaphores

Approach #2: Multiple Threads

- ▶ Very similar to approach #1 (multiple processes)
 - but, with threads instead of processes

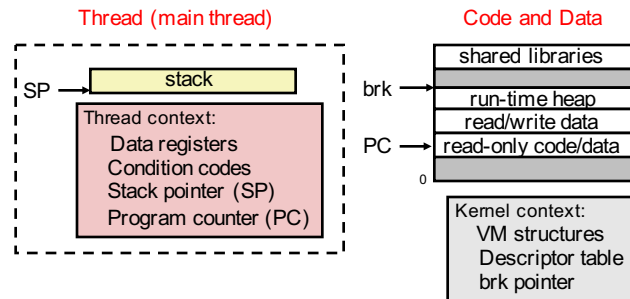
Traditional View of a Process

- ▶ Process = process context + code, data, and stack



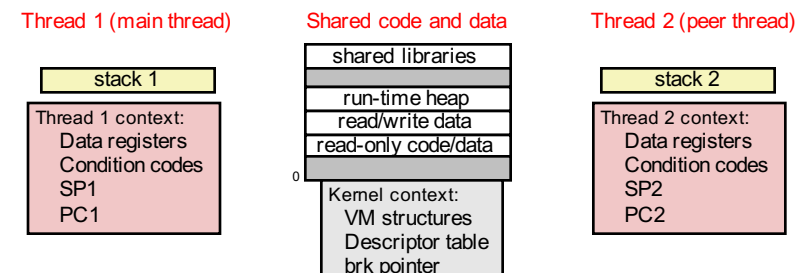
Alternate View of a Process

- ▶ Process = thread + code, data, and kernel context



A Process With Multiple Threads


- ▶ Multiple threads can be associated with a process
 - Each thread has its own logical control flow
 - Each thread shares the same code, data, and kernel context
 - Share common virtual address space (inc. stacks)
 - Each thread has its own thread id (TID)




Logical View of Threads

- Threads associated with process form a pool of peers
 - Unlike processes which form a tree hierarchy

Threads associated with a process

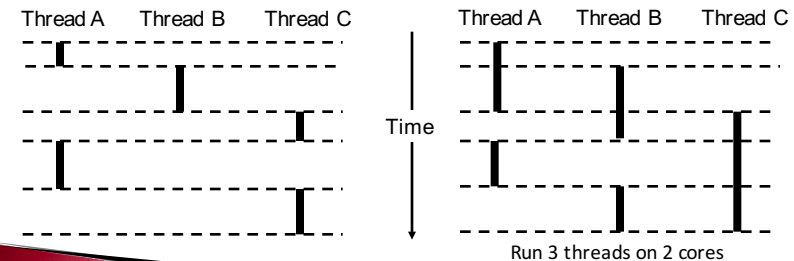


Process hierarchy



Thread Execution

- Single Core Processor
 - Simulate concurrency by time slicing
- Multi-Core Processor
 - Can have true concurrency




Logical Concurrency

- Two threads are (logically) concurrent if their flows overlap in time; otherwise, they are sequential

Examples:

- Concurrent: A & B, A&C
- Sequential: B & C



Threads vs. Processes

- How threads and processes are similar
 - Each has its own logical control flow
 - Each can run concurrently with others (possibly on different cores)
 - Each is context switched
- How threads and processes are different
 - Threads share code and some data
 - Processes (typically) do not
 - Threads are somewhat less expensive than processes
 - Process control (creating and reaping) is about twice as expensive as thread control

Posix Threads (Pthreads) Interface

- *Pthreads*: Standard interface for ~60 functions that manipulate threads from C programs
 - Creating and reaping threads
 - `pthread_create()`
 - `pthread_join()`
 - Determining your thread ID
 - `pthread_self()`
 - Terminating threads
 - `pthread_cancel()`
 - `pthread_exit()`
 - `exit()` [terminates all threads], `RET` [terminates current thread]
 - Synchronizing access to shared variables
 - `pthread_mutex_init`
 - `pthread_mutex_[un]lock`
 - `pthread_cond_init`
 - `pthread_cond_[timed]wait`

The Pthreads "hello, world" Program

```

/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"

void *thread(void *vargp);

int main() {
    pthread_t tid;

    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}

void *thread(void *vargp) {
    printf("Hello, world!\n");
    return NULL;
}

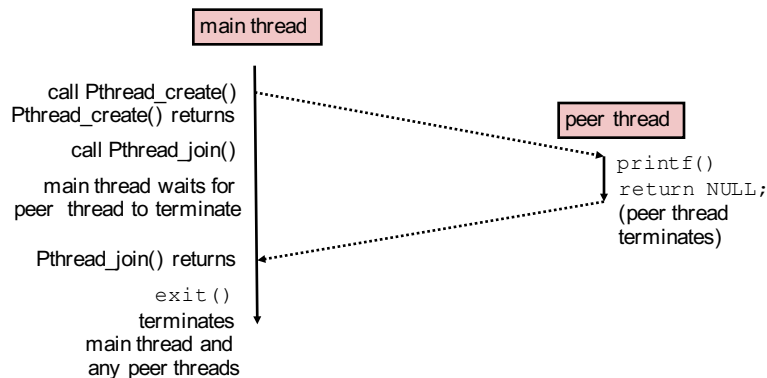
```

Thread attributes
(usually NULL)

Thread arguments
(void *p)

return value
(void **p)

Execution of Threaded "hello, world"



Thread-Based Concurrent Echo Server

```

int main(int argc, char **argv) {
    int port = atoi(argv[1]);
    struct sockaddr_in clientaddr;
    int clientlen = sizeof(clientaddr);
    pthread_t tid;

    int listenfd = Open_listenfd(port);
    while (1) {
        int *connfdp = Malloc(sizeof(int));
        *connfdp = Accept(listenfd,
                          (SA *)&clientaddr, &clientlen);
        Pthread_create(&tid, NULL, echo_thread, connfdp);
    }
}

```


- Spawn new thread for each client
- Pass it copy of connection file descriptor
- Note use of `Malloc()`!
 - Without corresponding `Free()`

Thread-Based Concurrent Server (cont)

```
/* thread routine */
void *echo_thread(void *vargp)
{
    int connfd = *((int *)vargp);
    Pthread_detach(pthread_self());
    Free(vargp);
    echo(connfd);
    Close(connfd);
    return NULL;
}
```

- Run thread in “detached” mode
 - Runs independently of other threads
 - Reaped when it terminates
- Free storage allocated to hold clientfd
 - “Producer-Consumer” model

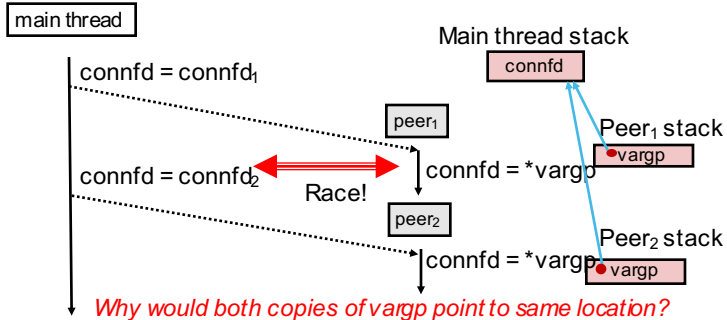
Threaded Execution Model



- Multiple threads within single process
- Some state between them
 - File descriptors

Potential Form of Unintended Sharing

```
while (1) {
    int connfd = Accept(listenfd, (SA *) &clientaddr, &clientlen);
    Pthread_create(&tid, NULL, echo_thread, (void *) &connfd);
}
}
```



Could this race occur?

Main

```
int i;
for (i = 0; i < 100; i++) {
    Pthread_create(&tid, NULL,
                  thread, &i);
}
```

Thread


```
void *thread(void *vargp)
{
    int i = *((int *)vargp);
    Pthread_detach(pthread_self());
    save_value(i);
    return NULL;
}
```

► Race Test


- If no race, then each thread would get different value of *i*
- Set of saved values would consist of one copy each of 0 through 99.

Experimental Results


No Race



Single core laptop



Multicore server



► The race can really happen!

Issues With Thread-Based Servers

- Must run “detached” to avoid memory leak.
 - At any point in time, a thread is either *joinable* or *detached*.
 - *Joinable* thread can be reaped and killed by other threads.
 - must be reaped (with `pthread_join`) to free memory resources.
 - *Detached* thread cannot be reaped or killed by other threads.
 - resources are automatically reaped on termination.
 - Default state is joinable.
 - use `pthread_detach(pthread_self())` to make detached.
- Must be careful to avoid unintended sharing.
 - For example, passing pointer to main thread's stack
`pthread_create(&tid, NULL, thread, (void *) &connfd);`
- All functions called by a thread must be *thread-safe*

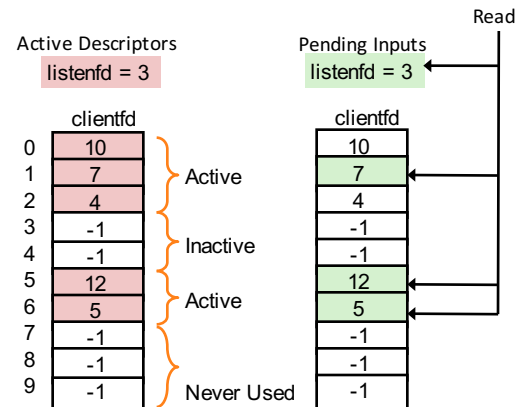
Pros and Cons of Thread-Based Designs

- + Easy to share data structures between threads
 - e.g., logging information, file cache.
- + Threads are more efficient than processes.
- – Unintentional sharing can lead to subtle, hard-to-reproduce errors!
 - The ease with which data can be shared is both the greatest strength and the greatest weakness of threads.
 - Hard to know which data shared & which private
 - Hard to detect by testing
 - Probability of bad race outcome very low
 - But nonzero!

Event-Based Concurrent Servers Using I/O Multiplexing

- Use library functions to construct scheduler within single process
- Server maintains set of active connections
 - Array of `connfd`'s
- Repeat:
 - Determine which connections have pending inputs
 - If `listenfd` has input, then accept connection
 - Add new `connfd` to array
 - Service all `connfd`'s with pending inputs
- Details in book

I/O Multiplexed Event Processing



Pros and Cons of I/O Multiplexing

- ▶ + One logical control flow.
- ▶ + Can single-step with a debugger.
- ▶ + No process or thread control overhead.
 - Design of choice for high-performance Web servers and search engines.
- ▶ – Significantly more complex to code than process- or thread-based designs.
- ▶ – Hard to provide fine-grained concurrency
 - E.g., our example will hang up with partial lines.
- ▶ – Cannot take advantage of multi-core
 - Single thread of control

Approaches to Concurrency

- ▶ Processes
 - Hard to share resources: Easy to avoid unintended sharing
 - High overhead in adding/removing clients
- ▶ Threads
 - Easy to share resources: Perhaps too easy
 - Medium overhead
 - Not much control over scheduling policies
 - Difficult to debug
 - Event orderings not repeatable
- ▶ I/O Multiplexing
 - Tedious and low level
 - Total control over scheduling
 - Very low overhead
 - Cannot create as fine grained a level of concurrency
 - Does not make use of multi-core

Synchronization

CS 341: Intro. to Computer
Architecture & Organization

Prof. Dorian Arnold

Why we need Synchronization

- ▶ Concurrent access to shared data can lead to inconsistencies
- ▶ Need synchronization amongst sharers to guarantee consistency

41

Producer/Consumer Example

```
while (true) {
    while (count == BUFFER_SIZE)
        ; // do nothing

    buffer [in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    count++;
}
```

Producer Code

42

Producer/Consumer Example

```
while (true) {
    while (count == 0)
        ; // do nothing

    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;
    /* consume the item nextConsumed */
}
```

Consumer Code

43

The Critical Section Problem

- `count++` could be implemented as

```
register1 = count
register1 = register1 + 1
count = register1
```

- `count--` could be implemented as

```
register2 = count
register2 = register2 - 1
count = register2
```

- Consider this execution interleaving with "count = 5" initially:

```
S0: producer execute register1 = count {register1 = 5}
S1: producer execute register1 = register1 + 1 {register1 = 6}
S2: consumer execute register2 = count {register2 = 5}
S3: consumer execute register2 = register2 - 1 {register2 = 4}
S4: producer execute count = register1 {count = 6}
S5: consumer execute count = register2 {count = 4}
```

44

Concurrency Terminology

- ▶ **Race condition** exists when output varies depending on access (write) order
- ▶ **Critical section**: section of code used by multiple processes (or threads) to update shared data
 - **Entry section**: request access to critical section
 - **Exit section**: notify departure from critical section
 - **Remainder section**: code outside critical section
- ▶ **lock**: abstraction/data structure used to access code in critical section

45

Critical Sections

```
do {
    entry section
    critical section
    exit section
    remainder section
} while( condition )
```

46

Solutions to Critical Section Problem

- ▶ **Mutual exclusion**: only one process can be executing code in a critical section at any point in time
- ▶ **Progress**: if no process is in critical section, process(es) wishing to enter the critical section must be allowed to do so
- ▶ **Bounded wait**: a process cannot be permanently blocked from entering a critical section

47

Enforcing Mutual Exclusion

- ▶ *Question*: How can we guarantee a safe trajectory?
- ▶ Answer: We must **synchronize** the execution of the threads so that they never have an unsafe trajectory.
 - i.e., need to guarantee **mutually exclusive access** to critical regions
- ▶ Classic solution:
 - Semaphores (Edsger Dijkstra)
- ▶ Other approaches (out of our scope)
 - Mutex and condition variables (Pthreads)
 - Monitors (Java)

Semaphores Operations

```

//initialization
binary semaphore (mutex lock) → S = 1;
counting semaphore → S = N;

wait(S)
{
    while ( S <= 0 ) ; ← spinlock
    S--;
}

signal( S )
{
    S++;
}

```

When might spinlocks be useful?

49

Entry and Exit using Semaphores

```

S = 1;

do {
    wait(S);

    // critical section

    signal(S);

    // remainder section

} while (TRUE);

```

50

Non-spin-lock Semaphores

```

typedef struct {
    int value;
    struct process *list;
} semaphore;

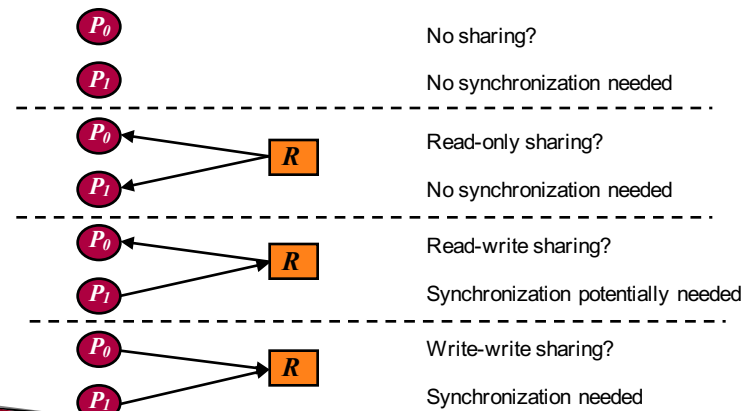
wait( semaphore *S ) {
    S->value--;
    if (S->value < 0) {
        S->list->add( current_process );
        block();
    }
}

signal( semaphore *S ) {
    S->value++;
    if (S->value <= 0) {
        S->list->remove( selected_process );
        wakeup( selected_process );
    }
}

```

51


Why Synchronization (Summary)



52

Synchronization Mechanisms

- ▶ Guarantee mutual exclusion when/where necessary
 - Higher-level primitives for programming convenience



Deadlock

- ▶ **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

- ▶ Let **S** and **Q** be two semaphores initialized to 1

P_0 wait (S); wait (Q); . . . signal (S); signal (Q);	P_1 wait (Q); wait (S); . . . signal (Q); signal (S);
--	--

54

Synchronization Issues

- ▶ Subtle synchronization errors
- ▶ Races
- ▶ Priority Inversion
- ▶ Deadlock